UNIVERSITY OF MINES AND TECHNOLOGY

FACULTY OF COMPUTING AND MATHEMATICAL SCIENCES

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

PROJECT REPORT ENTITLED

DYNAMIC CONTEXT-AWARE TRIES FOR KNOWLEDGE GRAPH-CONSTRAINED LARGE LANGUAGE MODEL GENERATION

BY

BERNARD KIRK ADJANOR KATAMANSO— FOE.41.008.088.22

AMONOR ERICA— FOE.41.008.030.22

JOSEPH OSEI NYARKO— FOE.41.008.113.22

JESSICA AFUA ETORNAM NSAFOAH— FOE.41.008.107.22

SUBMITTED IN PARTIAL FULFILMENT OF THE REQUIREMENT FOR THE AWARD OF THE DEGREE OF BACHELOR OF SCIENCE IN COMPUTER SCIENCE AND ENGINEERING

PROJECT SUPERVISOR

.....

Dr ERIC AFFUM

Tarkwa, Ghana

Month, Year

DECLARATION

We, BERNARD KIRK ADJANOR KATAMANSO, AMONOR ERICA, JOSEPH OSEI NYARKO and JESSICA AFUA ETORNAM NSAFOAH, declare that this project work is our own original work. It is being submitted for the degree of Bachelor of Science in Computer Science and Engineering at the University of Mines and Technology (UMaT), Tarkwa. It has not been submitted, either in whole or in part, for any degree or examination at any other university.

Signature of Student: _____

BERNARD KIRK ADJANOR KATAMANSO (FOE.41.008.088.22)

Signature of Student: _____

AMONOR ERICA (FOE.41.008.113.22)

Signature of Student: _____

JOSEPH OSEI NYARKO (FOE.41.008.107.22)

Signature of Student: _____

JESSICA AFUA ETORNAM NSAFOAH (FOE.41.008.030.22)

_____ day of _____ (year) _____

ABSTRACT

Large language models (LLMs) are capable of multi-hop reasoning but remain prone to hallucination. Knowledge graph-constrained decoding mitigates this by strictly restricting output tokens to valid paths within a knowledge graph (KG), ensuring absolute structural faithfulness unlike standard soft-grounding approaches (e.g., Retrieval-Augmented Generation). However, leading constrained paradigms—like Graph-Constrained Reasoning (GCR)—construct constraints solely from graph structure, ignoring the question's semantics and the evolving generation state.

This causes the *permissiveness problem*: at every decoding step, the allowed valid token set contains thousands of structurally accessible but contextually irrelevant paths. This flood of valid tokens forces the LLM to rely on the very parametric biases that constrained decoding was designed to suppress.

To resolve this, we propose Dynamic Context-Aware Tries (DCA-Trie), a novel framework that conditions the valid tokens jointly on the knowledge graph, the input question, and the partially generated reasoning chain. It upgrades the constraint oracle to an active semantic state $f(\mathcal{G}, E_q, q, y_{<t})$ by using a pretrained relevance scorer to filter candidate KG paths. Two architectures are developed: static semantic filtering applied at trie construction (DCA-Trie v1), and dynamic, step-wise bounding conditioned on the evolving generation (DCA-Trie v2).

Both variants are evaluated on the WebQSP and ComplexWebQuestions datasets over Freebase. Alongside standard metrics (Hits@1, F1), we introduce the Semantic Irrelevance Ratio (SIR)—a novel metric to uniquely isolate and quantify constraint permissiveness. To isolate and quantify constraint permissiveness, we introduce the Semantic Irrelevance Ratio (SIR), the first metric to measure the proportion of structurally valid but semantically irrelevant paths in the constraint set. Results demonstrate that DCA-Trie successfully eliminates semantically blind branches while preserving 100% structural faithfulness, charting a direct path toward more robust and efficient constrained reasoning.

DEDICATION

To our parents, who taught us that the truest things cannot be computed.

And to the machine, an elegant engine of probability, reminding us that predicting the word is not understanding the world.

ACKNOWLEDGMENTS

We extend our sincere gratitude to Dr. Eric Affum for his invaluable guidance, patience, and support throughout this project. His expertise in shaping our research direction has been instrumental in the completion of this dissertation.

We also wish to express our deep appreciation to our family, friends, and mentors. Their continuous encouragement, insight, and understanding provided the foundation we needed to see this work through to the end.

Finally, we acknowledge the broader academic community whose work has inspired and informed our research. This project stands on the shoulders of countless scholars, and we are grateful for their contributions to the field.

TABLE OF CONTENTS

DECLARATION **i**

ABSTRACT **ii**

ABSTRACT **ii**

DEDICATION **iii**

ACKNOWLEDGMENTS **iv**

TABLE OF CONTENTS **v**

LIST OF FIGURES **ix**

LIST OF TABLES **x**

LIST OF ABBREVIATIONS **xi**

CHAPTER 1 INTRODUCTION **1**

1.1 Background 1

1.2 Problem Statement 4

1.3 Aim and Objectives 6

1.4 Tools and Facilities Used 7

1.4.1 Software Tools 7

1.4.2 Facilities 9

1.5 Methods Used 9

1.6 Scope of Work 10

1.7 Organization of Work 11

CHAPTER 2 Literature Review **12**
2.1 Introduction 12
2.2 Hallucination in Large Language Models: Taxonomy, Structural Causes, and the Scaling Paradox 13
2.2.1 Defining Hallucination: A Formal Taxonomy for Knowledge-Intensive Generation 13
2.2.2 The Autoregressive Generation Process as the Structural Root of Hallucination 14
2.2.3 The Scaling Paradox: Why Larger Models Do Not Solve Hallucination 14
2.3 Chain-of-Thought Reasoning, Self-Consistency, and the Faithfulness Ceiling of Unconstrained Multi-Step Generation 16
2.3.1 Chain-of-Thought Prompting: Eliciting Step-by-Step Reasoning from LLMs 16
2.3.2 Self-Consistency and Tree-of-Thought: Improving Reliability Without Improving Faithfulness 17
2.3.3 The Faithfulness Ceiling: The Mathematical Limits of Prompting 18
2.4 Knowledge Graph Question Answering: Structured Reasoning, Benchmark Design, and the Exponential Complexity of Multi-Hop Inference 19
2.4.1 Knowledge Graphs as Structured Repositories of Verified Facts: Formal Definition and Properties 19
Concrete Example: A Knowledge Graph Subgraph from Freebase 20
2.4.2 WebQSP and ComplexWebQuestions: Benchmark Architecture, Evaluation Protocol, and Hop-Depth Stratification 21
2.4.3 Multi-Hop Reasoning Complexity: Exponential Path Growth and the Permissiveness Problem Quantified 22
2.5 Constrained Decoding for Faithful Text Generation: From Logit Masking to Knowledge Graph Oracles 24
2.5.1 The Logit-Masking Paradigm: Hard Constraints on Autoregressive Vocabulary Selection 24
2.5.2 Grammar-Constrained Decoding: GCD and Outlines as Foundational Format Constraints 25

2.5.3 Graph-Constrained Reasoning (GCR): Structural Faithfulness via KG-Trie Oracles 26
2.5.4 Decoding on Graphs (DoG): Step-Wise Topological Expansion and Its Persistent Limitations 27
2.6 Beam Search Decoding and Prefix-Tree Constraint Integration: Mechanics, Efficiency, and the Step-Wise Constraint Problem 29
2.6.1 Beam Search: Approximate Global Optimization Over Autoregressive Sequences 29
2.6.2 Prefix Trees as Constraint Oracles: Efficiency, Precomputation, and the Dynamic Rebuilding Cost 30
2.6.3 The Step-Wise Constraint Problem: When the Constraint Must Know What Has Been Generated 31
2.7 Retrieval-Augmented Generation: Architectural Variants, KG-Augmented Approaches, and the Hard Limit of Soft Grounding 32
2.7.1 The RAG Paradigm: Dense Retrieval, Fusion-in-Decoder, and the Soft Grounding Ceiling 32
2.7.2 KG-Augmented Retrieval: Structured Context Without Structural Guarantees 33
2.7.3 Advanced RAG Variants: Self-RAG, FLARE, IRCoT, and GraphRAG 33
2.7.4 The Formal Impossibility of Structural Faithfulness Under Soft Grounding 34
2.8 Semantic Relevance Scoring with Sentence Transformer Embeddings: Training Objectives, Architecture Selection, and Application to Constraint Oracle Construction 35
2.8.1 Sentence-BERT: Training Objective, Architecture, and the Semantic Similarity Property 35
2.8.2 Encoder Architecture Selection: all-MiniLM-L6-v2 vs Alternatives 36
2.8.3 Cosine Similarity as a Hard Constraint Signal: From Soft Retrieval to Binary Admission 37
2.9 Related Work 38
2.9.1 KG-Constrained Decoding: GCR, DoG, and the Unresolved Permissiveness Problem 38

2.9.2 Hallucination Mitigation and KGQA: The Soft-to-Hard Spectrum 39
2.9.3 Semantic Similarity for Graph Reasoning: The Retrieval-Constraint Boundary 39
2.9.4 Concurrent and Subsequent Work: Positioning DCA-Trie 40
Differentiation of DCA-Trie 40
2.10 Synthesis of the Literature: Three Unresolved Gaps and the DCA-Trie Contribution 41
2.10.1 Three Unresolved Gaps at the Intersection of Structural Faithfulness, Semantic Relevance, and Constraint Measurability 41
2.10.2 DCA-Trie as the Unified Response: Gap-to-Objective Mapping 42

REFERENCES **44**

LIST OF FIGURES

Figure	Title	Page
2.1	Relationship between model scale and truthfulness on the TruthfulQA benchmark.	15
2.2	Example knowledge graph subgraph from Freebase	21
2.3	Hop-depth distributions for WebQSP and ComplexWebQuestions.	22
2.4	Exponential growth of structurally valid KG paths with hop depth.	23
2.5	The logit-masking mechanism. Token selection operates exclusively over the valid, non-zero subset defined by the constraint oracle mask.	25
2.6	Grammar-Constrained Decoding Algorithm.	25
2.7	The GCR KG-Trie illustrating the permissiveness problem.	27
2.8	Evolution of the constraint valid set across decoding steps.	29
2.9	Interaction between prefix-tree constraint lookup and beam search expansion.	31
2.10	Architectural comparison of standard Retrieval-Augmented Generation (left) and KG-constrained decoding (right).	34
2.11	End-to-end semantic scoring and binary admission pipeline.	38

LIST OF TABLES

Table	Title	Page
1.1	Implemented Software Tools Evaluation	8
2.1	Comparison of chain-of-thought family methods and KG-constrained decoding.	18
2.2	Comparison of sentence encoder architectures.	36
2.3	Primary related work comparison.	39
2.4	Architectural comparison of RouterKGQA and DCA-Trie.	40
2.5	Consolidated gap matrix for KG-constrained decoding frameworks.	42
2.6	DCA-Trie contribution mapping.	43

LIST OF ABBREVIATIONS

DCA-Trie	Dynamic Context-Aware Trie
LLM	Large Language Model
KGQA	Knowledge Graph Question Answering
RAG	Retrieval-Augmented Generation
KG	Knowledge Graph
GCR	Graph-Constrained Reasoning
CWQ	ComplexWebQuestions
RLHF	Reinforcement Learning from Human Feedback
CoT	Chain-of-Thought
DoG	Decoding on Graphs
SIR	Semantic Irrelevance Ratio
GCD	Grammar-Constrained Decoding
QA	Question Answering
NLP	Natural Language Processing
AI	Artificial Intelligence
BFS	Breadth-First Search

CHAPTER 1

INTRODUCTION

1.1 Background

Large language models (LLMs) have emerged as one of the most consequential advances in artificial intelligence, demonstrating capabilities across language understanding, generation, and multistep reasoning that were considered beyond reach only a decade ago (Brown *et al.*, 2020; OpenAI, 2024). Models such as the GPT series (Brown *et al.*, 2020; Achiam *et al.*, 2023) and the Llama family (Touvron *et al.*, 2023; Grattafiori *et al.*, 2024) achieve strong performance on a broad spectrum of language tasks by learning statistical regularities over vast corpora of text. The scale and breadth of this training enables these models to exhibit apparent knowledge across diverse domains, from history and science to legal reasoning and medical information, without any explicit representation of factual knowledge in their architecture.

Despite these achievements, a fundamental and persistent failure mode undermines the reliability of LLMs in knowledge-intensive settings: hallucination. Ji *et al.* (2023) define hallucination as the generation of content that is fluent and coherent but factually unsupported or contradicted by verifiable sources. Huang *et al.* (2023) provide a comprehensive taxonomy distinguishing intrinsic hallucination, where the output contradicts a provided source, from extrinsic hallucination, where the output cannot be verified against any source. Critically, hallucination is not a rare edge case. Luo *et al.* (2024) report that state-of-the-art retrieval-augmented approaches produce hallucinated reasoning steps in approximately one-third of cases even when granted access to a supporting knowledge graph. This failure rate renders unconstrained LLMs unsuitable for high-stakes knowledge-intensive tasks without additional mechanisms to enforce factual grounding. The structural inevitability of hallucination has been formally established in recent theoretical work. Xu *et al.* (2025) prove that for any computable ground truth function, no LLM training procedure can achieve zero hallucination probability across all inputs–hallucination is an inherent property of computable learning systems rather than a transient training artifact. This "inevitability" result underscores that mitigation strategies relying solely on improved training, scaling, or prompting cannot fully

eliminate hallucination; external structural constraints that ground generation in verifiable sources are theoretically necessary for guaranteed faithfulness. Knowledge graph-constrained decoding provides exactly such an external mechanism, making hallucination mathematically impossible rather than merely improbable.

The structural cause of hallucination is well understood. Autoregressive LLMs generate text by conditioning each token on all previously generated tokens, producing a probability distribution over the vocabulary at each step: $P(y_t \mid y_{<t}, x)$, where x is the input and $y_{<t}$ is the partially generated sequence (Brown *et al.*, 2020). This process is effective at producing locally coherent text but provides no mechanism for grounding each generated claim in a verifiable external source. An error introduced at step t propagates forward as the conditioning context for all subsequent tokens, producing confident and grammatically well-formed outputs whose factual incorrectness has no surface signal (Huang *et al.*, 2023; Ji *et al.*, 2023).

Knowledge graph question answering (KGQA) represents one of the most demanding contexts in which this failure mode manifests. A knowledge graph (KG) is a structured, human-curated database of verified facts encoded as (entity, relation, entity) triplets, for instance, (Marie Curie, award received, Nobel Prize in Physics) or (Paris, capital of, France), forming a large-scale graph over which multi-hop reasoning paths can be traced (Bollacker *et al.*, 2008). The Freebase knowledge graph, which provides the underlying resource for the standard KGQA benchmarks WebQSP (Yih *et al.*, 2016) and ComplexWebQuestions (Talmor and Berant, 2018), encodes tens of millions of such triplets covering a wide range of domains and relation types. Answering a multi-hop question such as 'What is the nationality of the director of the film in which the Blue Hawaii actor starred?' requires traversing a chain of three or more KG triplets in the correct order, a task for which an unconstrained LLM, relying entirely on parametric memory, may produce a plausible-sounding but factually incorrect chain that is indistinguishable in fluency from a correct one (Lan *et al.*, 2022; He *et al.*, 2021a).

The dominant paradigm for addressing this problem is retrieval-augmented generation (RAG), in which relevant passages or subgraphs are retrieved from an external source and provided to the LLM as additional context (Lewis *et al.*, 2020; Izacard and Grave, 2021). While RAG reduces hallucination by providing relevant source material, it does not prevent the LLM from ignoring that material or generating claims that go beyond it. The LLM remains free to

produce any token at any step; the retrieved context is a soft influence on its output, not a hard constraint on what it may generate (Asai *et al.*, 2024; Agrawal *et al.*, 2024b).

KG-constrained decoding addresses this limitation directly. Rather than providing the KG as a soft contextual influence, constrained decoding frameworks interpose a constraint oracle between the LLM's logit distribution and the token selection mechanism. At each generation step, the oracle computes a binary mask over the vocabulary, setting to zero the logit of any token that would not correspond to a valid continuation of a KG path. The masking formulation, introduced by Geng *et al.* (2023) in the context of grammar-constrained decoding, takes the form $\tilde{\alpha}_t = m_t \odot \alpha_t$, where α_t is the raw logit vector, $m_t \in \{0, 1\}^{|V|}$ is the validity mask, and $\odot$ denotes element-wise product. The generation process thus selects from a restricted set of valid tokens at every step, making structural hallucination, the generation of facts that do not exist in the KG, mathematically impossible rather than merely improbable (Willard and Louf, 2023; Geng *et al.*, 2023).

The leading implementation of this paradigm is Graph-Constrained Reasoning (GCR) (Luo *et al.*, 2024), which achieves 92.6 Hits@1 on WebQSP and 75.8 Hits@1 on CWQ with verified 100% structural faithfulness. GCR constructs a KG-Trie, a prefix tree (Fredkin, 1960) over all KG paths within L hops of the question entities, obtained by breadth-first search, and uses it as the constraint oracle during beam-search decoding. Because every token in the generated reasoning chain must correspond to a valid prefix in the trie, the output is guaranteed to consist entirely of verified KG facts. This represents a qualitative improvement over both unconstrained LLMs and soft RAG approaches for the specific property of structural faithfulness.

However, a significant limitation remains. The KG-Trie in GCR is computed once, before generation begins, from the graph structure and the question entity set alone. It carries no information about the question semantics, the direction in which reasoning is proceeding, or which entities and relations have already been committed to in the partial generation $y_{<t}$. Formally, the valid token set W_{val} is a function $f(G, E_q)$, independent of q and $y_{<t}$, and remains identical across every generation step. For a question about a specific narrow fact, this trie may contain hundreds or thousands of structurally valid but semantically irrelevant paths, forcing the LLM to discriminate among them using its own parametric knowledge: precisely the hallucination-prone mechanism that constrained decoding was designed to

suppress (Li *et al.*, 2024; Luo *et al.*, 2024).

Decoding on Graphs (DoG) (Li *et al.*, 2024) introduces dynamic step-wise expansion, updating the constraint trie after each committed entity to include KG neighbors of the current node. This topological dynamism ensures the constraint set reflects the reasoning direction taken, addressing GCR's structural rigidity. DoG also employs question semantics, but crucially, only for *ranking* candidate paths during beam search–the valid token set itself remains determined solely by graph connectivity, admitting every neighbor regardless of relevance to the question or partial generation. The permissiveness problem therefore persists: thousands of structurally valid but semantically irrelevant paths remain in the valid set at each step, forcing the LLM to discriminate among them using parametric knowledge. Furthermore, DoG sacrifices the computational efficiency of pre-computed tries, rebuilding the constraint structure at each decoding step. These limitations motivate the DCA-Trie framework, which applies semantic filtering directly to the constraint oracle, ensuring that the valid token set at each step is both structurally faithful and semantically relevant.

1.2 Problem Statement

Existing KG-constrained decoding frameworks, including GCR (Luo *et al.*, 2024) and DoG (Li *et al.*, 2024), construct their constraint oracles as functions $f(G, E_q)$ of the knowledge graph and question entities, evaluated prior to or independently of the sequential generation process. This design has been shown to guarantee structural faithfulness (every generated token corresponds to a verified KG triplet), but provides no mechanism for adapting the constraint to the semantic content of the question or the direction taken by the partial generation (Geng *et al.*, 2023; Willard and Louf, 2023).

The consequence is a fundamental conflation of two distinct requirements. Structural validity requires that every generated token correspond to a valid KG path prefix, a requirement all existing frameworks satisfy. Contextual relevance requires that the constraint set reflect the semantic focus of the reasoning at the current step, a requirement no existing framework satisfies. As a result, existing constraint oracles are simultaneously too broad, admitting structurally valid but semantically irrelevant paths that force the model to rely on parametric discrimination, and too rigid, failing to narrow in response to the evolving reasoning direction as generation progresses (Li *et al.*, 2024; Luo *et al.*, 2024).

This project addresses the following specific problem: current KG-constrained decoding frameworks do not condition the valid token constraint on the input question semantics or the partial generation, leaving the LLM responsible for semantic discrimination among a large set of valid-but-irrelevant paths at every generation step. This residual reliance on parametric knowledge partially re-introduces the hallucination risk that constrained decoding was designed to eliminate, and is most pronounced in multi-hop settings where the number of valid paths grows exponentially with hop depth (Lan *et al.*, 2022; He *et al.*, 2021a; Talmor and Berant, 2018).

Very recent work has independently addressed the efficiency-faithfulness trade-off in KGQA. RouterKGQA (Yuan *et al.*, 2026) proposes routing between specialized and general models with semantic filtering at the answer level, achieving strong performance with minimal inference cost. While RouterKGQA and DCA-Trie both employ semantic relevance scoring, they operate at fundamentally different architectural stages: RouterKGQA filters *answers* after generation, whereas DCA-Trie filters *tokens* during generation. This distinction is critical: DCA-Trie's constraint oracle ensures the LLM never sees semantically irrelevant tokens, reducing cognitive load and enabling deterministic faithfulness guarantees that post-hoc filtering cannot provide. Furthermore, DCA-Trie introduces the Semantic Irrelevance Ratio (SIR; Chapter **??**), the first metric to quantify constraint permissiveness independently of downstream accuracy–an evaluation dimension not addressed by routing-based approaches.

To address these limitations, this project proposes DCA-Trie, upgrading the constraint oracle from $W_{\text{val}}(t) = f(G, E_q)$ to $W_{\text{val}}(t) = f(G, E_q, q, y_{<t})$. The upgraded framework must simultaneously satisfy three conditions: (i) structural faithfulness is preserved, every token admitted at each step corresponds to a valid KG path, as verified by exhaustive triplet checking against Freebase; (ii) semantic irrelevance is reduced, the proportion of the valid set consisting of paths irrelevant to the question decreases relative to GCR and DoG, as measured by the Semantic Irrelevance Ratio (SIR) introduced in this work; and (iii) answer recall is maintained: no gold answer path is erroneously excluded by the relevance filter, with false negative rate held below 5% on the WebQSP validation set.

1.3 Aim and Objectives

This project aims to design, implement, and evaluate a Dynamic Context-Aware Trie (DCA-Trie) framework for knowledge graph-constrained large language model generation, in which the valid token constraint at each decoding step is conditioned jointly on the knowledge graph structure, the input question semantics, and the partially generated reasoning chain, achieving improved answer accuracy on multi-hop KGQA benchmarks while maintaining 100% structural faithfulness to the underlying knowledge graph.

To achieve this aim, the following objectives have been set:

i. To characterize the permissiveness problem in existing KG-constrained decoding frameworks by defining and measuring the Semantic Irrelevance Ratio (SIR) of GCR's static KG-Tries on the WebQSP and CWQ benchmarks, stratified by hop depth.

ii. To design a semantic relevance scoring mechanism using pretrained sentence transformer embeddings (Reimers and Gurevych, 2019) that can filter candidate KG paths by joint relevance to the input question and partial generation at inference time, without task-specific fine-tuning.

iii. To implement a static semantic filtering variant (DCA-Trie v1) that applies the relevance scorer at KG-Trie construction time, reducing trie size before generation begins while maintaining a false negative rate below 5% on gold answer paths on the WebQSP validation set.

iv. To implement a step-wise dynamic expansion variant (DCA-Trie v2) in which the trie is updated after each committed triplet using the relevance scorer conditioned jointly on the question and the partial generation $y_{<t}$, directly operationalizing the $W_{\text{val}}(t) = f(G, E_q, q, y_{<t})$ formulation.

v. To evaluate DCA-Trie v1 and DCA-Trie v2 against the GCR baseline and an unconstrained chain-of-thought (CoT) baseline on WebQSP and CWQ, measuring Hits@1, F1, structural faithfulness, SIR, and average trie size, with results stratified by question hop depth.

vi. To demonstrate the practical applicability of DCA-Trie by building an experimental interactive prototype system (web-based or command-line QA interface) for research

demonstration, allowing users to input complex natural language questions and receive factually grounded reasoning chains and answers, thereby showcasing the framework's real-world value beyond benchmark metrics.

1.4 Tools and Facilities Used

1.4.1 Software Tools

Table 1.1 presents an evaluation of the software tools implemented in this project, their role within the DCA-Trie framework, and the rationale for their selection.

Table 1.1 Implemented Software Tools Evaluation

Tool	Use in Project	Rationale
Python / HuggingFace Transformers	Model inference pipeline; loading and running GCR's Llama-3.1-8B checkpoint; integrating the relevance scorer	Provides reliable model loading, tokenization, and generation utilities, and enables direct use of the released GCR checkpoint.
GCR Framework	Primary baseline; KG-Trie construction, graph-constrained decoding, and result reproduction on WebQSP and CWQ	Provides a strong baseline and supports the trie-based constrained decoding mechanism that DCA-Trie extends.
Sentence-transformers / all-MiniLM-L6-v2	Semantic relevance scoring; computing cosine similarity between question-context and candidate KG path embeddings	Provides fast semantic embeddings without extra fine-tuning and supports efficient iterative decoding.
PyTorch	Deep learning framework underpinning all model inference and custom decoding logic	Supports implementation and testing of custom decoding logic and integrates smoothly with HuggingFace.
Freebase / WebQSP and CWQ Datasets	Knowledge graph source and KGQA benchmarks for evaluation	Provides standard datasets for this research area and enables direct comparison with prior work.
Google Colab Pro (A100, 40 GB)	GPU environment for all inference experiments; Llama-3.1-8B constrained decoding with beam search	Provides enough GPU memory for stable experiments and removes the need for local high-end hardware.
Visual Studio Code + Git / GitHub	Development environment and version control for all four team members	Supports day-to-day coding, team collaboration, version tracking, and reproducibility throughout the project.

1.4.2 Facilities

The following facilities were utilized in the execution of this project:

i. Google Colab Pro (A100, 40 GB VRAM) for all model inference and constrained decoding experiments.

ii. University library and online academic databases (ACM Digital Library, arXiv, Google Scholar, ACL Anthology) for literature review and citation verification.

iii. University internet infrastructure for cloud connectivity, dataset access, and retrieval of model checkpoints from the HuggingFace model hub.

iv. Shared team repository hosted on GitHub for version control, experiment logging, and reproducibility of all reported results.

1.5 Methods Used

The project employs the following methods:

i. Systematic review and critical analysis of the academic literature on constrained decoding, knowledge graph question answering, hallucination in LLMs, and retrieval-augmented generation.

ii. Formal gap analysis of the four foundational papers, GCR (Luo *et al.*, 2024), DoG (Li *et al.*, 2024), GCD (Geng *et al.*, 2023), and Outlines (Willard and Louf, 2023), to identify the specific limitations that motivate the DCA-Trie design, including the permissiveness problem and the absence of generation-state feedback in the constraint oracle.

iii. Formal specification of the DCA-Trie framework, including definition of the Semantic Irrelevance Ratio metric, the relevance scoring mechanism using sentence transformer embeddings (Reimers and Gurevych, 2019), the static filtering procedure for v1, and the step-wise dynamic expansion rule for v2.

iv. Reproduction of the GCR baseline on the WebQSP benchmark using GCR's publicly released Llama-3.1-8B checkpoint to establish a verified starting point within 1–2% of published Hits@1.

v. Implementation of DCA-Trie v1 and v2 as extensions to GCR's inference pipeline, with relevance threshold selected on a 100-question held-out validation subset of WebQSP to minimize false negative rate on gold answer paths.

vi. Quantitative evaluation of all system configurations on WebQSP and CWQ using Hits@1, F1, faithfulness rate, SIR, and average trie size per step, stratified by question hop depth following the evaluation protocol of Lan *et al.* (2022).

vii. Ablation studies examining the sensitivity of DCA-Trie's performance to the relevance threshold τ, the sentence encoder architecture, and beam size k, isolating the contribution of each mechanism to the overall result.

viii. Qualitative error analysis identifying failure cases in which DCA-Trie's semantic filter erroneously excludes a gold answer path, supporting honest characterization of the framework's limitations and providing concrete directions for future work.

1.6 Scope of Work

This project designs and evaluates a Dynamic Context-Aware Trie framework for KG-constrained LLM decoding, with the specific aim of addressing the permissiveness problem identified in GCR and DoG. The framework is developed and evaluated in the context of multi-hop KGQA over Freebase-grounded benchmarks (WebQSP and CWQ), which are the standard evaluation settings in the KG-constrained decoding literature. All experiments use GCR's publicly released Llama-3.1-8B checkpoint; no new model training or fine-tuning is performed. The semantic relevance scorer uses the all-MiniLM-L6-v2 sentence transformer without domain-specific adaptation. We emphasize that all experiments use GCR's publicly released Llama-3.1-8B checkpoint (Luo *et al.*, 2025), which includes their fine-tuned KG-specialized parameters. No additional model training or fine-tuning is performed in this work; our contribution is strictly the constraint oracle architecture and semantic filtering mechanism.

The scope is explicitly bounded as follows. Generalization of DCA-Trie to knowledge graphs other than Freebase, or to non-KGQA tasks such as entity linking or structured information extraction, is identified as future work and is not evaluated in this dissertation. Formal proofs of the faithfulness guarantee under dynamic pruning are beyond scope; empirical

measurement of the gold answer false negative rate is provided as evidence. The confidence-conditioned constraint tightening mechanism, in which constraint granularity is modulated by the LLM's token-level entropy, is described as a theoretical extension but is not implemented. Production-scale deployment, latency optimization for online serving, and integration with proprietary enterprise knowledge graphs are engineering extensions beyond the research scope of this project.

The primary contribution of this work is the DCA-Trie framework and the Semantic Irrelevance Ratio metric. Secondary contributions are the empirical characterization of constraint permissiveness in existing frameworks and the ablation analysis isolating the contribution of semantic filtering from dynamic expansion. All code, experimental configurations, and the curated evaluation datasets will be made publicly available to support reproducibility and future extension.

1.7 Organization of Work

Chapter 1 introduces the study by identifying the core problem, stating the project objectives, and defining the intended DCA-Trie solution, thereby establishing the motivation, scope, and overall direction of the research. Chapter 2 presents the literature review by examining prior work on hallucination in LLMs, KGQA, constrained decoding, and related frameworks, and it identifies key limitations in existing approaches to justify the need for a context-aware constraint mechanism. Chapter 3 combines problem formulation and methodology, defining the upgraded constraint setup, research questions, evaluation criteria, and implementation strategy for DCA-Trie v1 and v2 within the GCR pipeline (Luo *et al.*, 2024). Chapter 4 presents results and discussion, reporting quantitative evaluation across WebQSP and CWQ, analysing faithfulness and semantic irrelevance trends, and interpreting the findings through ablation and error analysis. Chapter 5 concludes the dissertation by summarizing the main contributions, highlighting practical implications, and presenting recommendations for future work.

CHAPTER 2

LITERATURE REVIEW

2.1 Introduction

Large Language Models (LLMs) have fundamentally transformed natural language processing by demonstrating unprecedented fluency and reasoning capabilities. However, their deployment in knowledge-intensive domains, where factual reliability is a strict prerequisite rather than a desirable trait, remains severely constrained. In contexts such as automated reasoning, medical diagnostics, and structured data querying, AI systems must do more than approximate plausible answers; they must ground their outputs in deterministically verifiable knowledge.

The foundational challenge hindering this adoption is hallucination: the generation of fluent but factually incorrect or unsupported claims. Because these models generate text autoregressively over an unconstrained vocabulary, they lack a built-in mechanism to ensure structural faithfulness to an external factual source. While current mitigation strategies, including chain-of-thought prompting and soft retrieval augmentation, significantly improve empirical accuracy, they fundamentally fail to provide hard, mathematical guarantees against hallucinated reasoning steps.

To address these challenges, this chapter examines existing research on these vulnerabilities and explores the progression of constrained decoding architectures designed to solve them. By doing so, it establishes the theoretical foundations for the Dynamic Context-Aware Trie (DCA-Trie). The review is structured around nine thematic pillars. First, Sections 2.2 and 2.3 formalize the problem by diagnosing hallucination's structural causes and exposing the faithfulness ceiling of unconstrained multi-step reasoning. Section 2.4 establishes the technical domain in which DCA-Trie operates: knowledge graph question answering and the mechanics of multi-hop inference. Sections 2.5 and 2.6 form the methodological core, tracing the evolution of constrained decoding paradigms, from foundational logit masking to leading graph-specific implementations like GCR and DoG, and analyzing their interaction with beam search. Section 2.7 examines retrieval-augmented generation to formally prove why soft grounding falls short of structural faithfulness, while Section 2.8 justifies the sentence

transformer embeddings used as the semantic relevance scorer at the heart of DCA-Trie. Finally, Sections 2.9 and 2.10 position this work against specific prior art and synthesize three unresolved research gaps that this project directly answers.

2.2 Hallucination in Large Language Models: Taxonomy, Structural Causes, and the Scaling Paradox

This section explores the root cause of the problem DCA-Trie sets out to solve. Rather than treating hallucination as a simple bug that can be patched, we examine it as a fundamental consequence of how autoregressive language models operate, meaning its elimination requires direct external constraints. To conclude, the section addresses a common counterargument, explaining why simply scaling up to larger models does not make these constraint mechanisms unnecessary.

2.2.1 Defining Hallucination: A Formal Taxonomy for Knowledge-Intensive Generation

Ji *et al.* (2023) define hallucination in natural language generation as the production of content that is fluent and internally coherent yet factually unsupported by, or in outright contradiction with, any verifiable source. Huang *et al.* (2023) refine this definition into a two-part taxonomy that is foundational to the present work. *Intrinsic hallucination* occurs when the model output directly contradicts material that is explicitly present in the provided source context. *Extrinsic hallucination* occurs when the output contains claims that cannot be verified against any provided or accessible external source, even if those claims are not outright contradictions.

When applied to the specific domain of knowledge graph question answering (KGQA), this two-part taxonomy translates into highly concrete, operational definitions. Intrinsic hallucination in KG-constrained generation means generating a triplet (e, r, e') that does not exist as a verified fact in the knowledge graph $\mathcal{G}$. Extrinsic hallucination means generating a multi-hop reasoning chain that cannot be traced to any verified path in $\mathcal{G}$, regardless of whether each individual generated token seems plausible in isolation. Both forms are unacceptable for safety-critical or knowledge-intensive applications, where factual reliability is a hard requirement rather than a desirable property. The severity of the problem is underscored by the observation in Luo *et al.* (2024) that even leading knowledge-graph-enhanced reasoning

methods, such as the RoG framework, still produce hallucinated reasoning steps in approximately 33% of cases. This single finding constitutes the strongest empirical justification for hard constraint mechanisms over soft grounding approaches.

2.2.2 The Autoregressive Generation Process as the Structural Root of Hallucination

Crucially, hallucination should not be viewed as a mere technical anomaly or a transient issue that can be easily trained away; rather, it is a fundamental architectural consequence of the autoregressive generation mechanism itself. At each generation step t, an autoregressive LLM computes a probability distribution over the full vocabulary conditioned on the input prompt x and all previously generated tokens $y_{<t} = (y_1, \ldots, y_{t-1})$:

$$P(y_t \mid y_{<t},\, x) \tag{2.1}$$

Three structural observations follow directly from this formulation (Brown *et al.*, 2020; Huang *et al.*, 2023; Ji *et al.*, 2023). First, no term in Equation (2.1) encodes a factual verification check: the distribution is computed entirely from learned weights and the conditioning context, with no mechanism for querying an external ground-truth source. Second, because $y_{<t}$ conditions all future steps, an error introduced at step $t-1$ propagates forward as part of the conditioning context for every subsequent token, producing grammatically fluent continuations whose factual incorrectness has no surface signal distinguishable from a correct chain. Third, the distribution is defined over the full vocabulary $\mathcal{V}$; every token has strictly positive probability at every step. The consequence is that structural hallucination is not a low-probability edge case; it is an ever-present possibility at every decoding step, and the fluency of the generated text provides no diagnostic signal that an error has occurred.

2.2.3 The Scaling Paradox: Why Larger Models Do Not Solve Hallucination

A natural and frequently encountered objection to constraint-based approaches is that sufficiently large models, or those trained with reinforcement learning from human feedback (RLHF), should be factually reliable enough to render hard constraints unnecessary. The evidence does not support this objection, and establishing this point is essential for a principled motivation of DCA-Trie. Lin *et al.* (2022) introduce TruthfulQA, a benchmark designed to measure whether models generate truthful answers by targeting questions that elicit common

human misconceptions. Their central finding is that larger models do not systematically outperform smaller models on truthfulness and that, for certain question categories, truthfulness accuracy actually decreases with scale, a phenomenon termed *inverse scaling*. This occurs because larger models natively become better at memorising and reproducing the false but highly prevalent misconceptions present in their training data.

As illustrated schematically in Figure 2.1, Truthfulness accuracy explicitly exhibits inverse scaling across base language models such as the GPT-3 and GPT-Neo families. Unconstrained language modeling performance fails to improve monotonically with parameter count and instead regresses robustly into the inverse-scaling region. By contrast, models explicitly fine-tuned on QA tasks (such as UnifiedQA) exhibit positive scaling, demonstrating that scaling standard autoregressive imitation alone inherently penalizes factual reasoning.

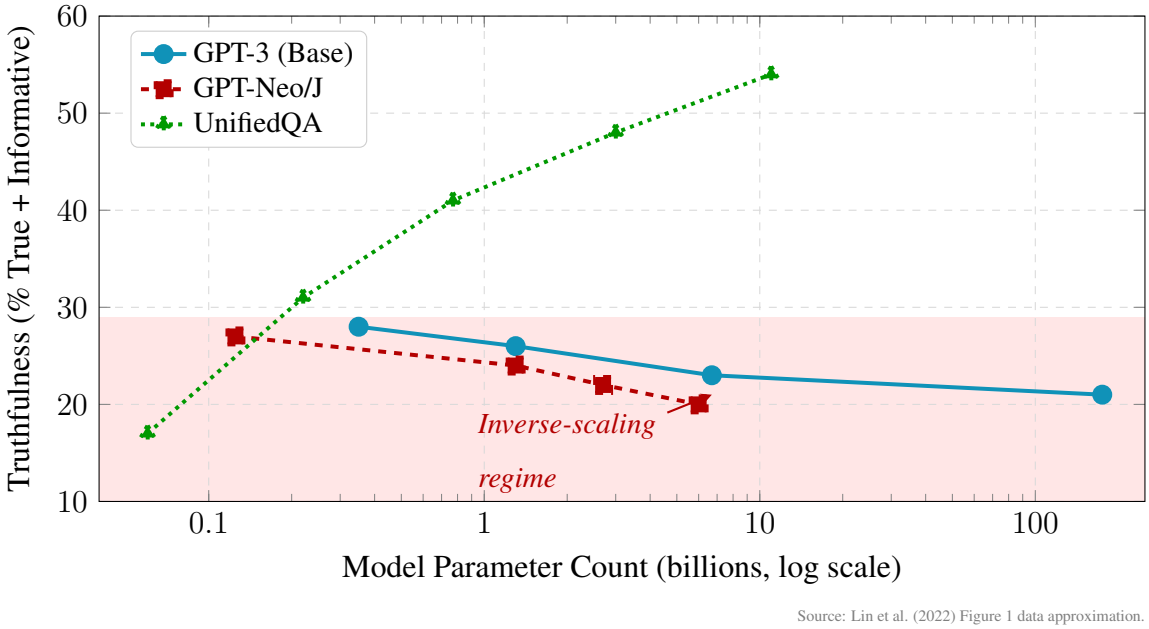

Figure 2.1 Relationship between model scale and truthfulness on the TruthfulQA benchmark.

The sycophancy problem identified by Perez *et al.* (2023) compounds this picture. RLHF-tuned models are trained to produce outputs aligned with user preferences; since users tend to rate confident and fluent responses favorably, larger RLHF-tuned models are more likely to generate confident-sounding but factually incorrect statements that match expected answers rather than true ones. Confidence and accuracy are thus decoupled by the very training procedure that was supposed to improve model reliability. Kandpal *et al.* (2023) identify a further limitation through the memorisation-generalisation tension: larger models memorize

more training data and therefore perform better on factual recall tasks involving frequently seen entity combinations, but they do not improve on compositional multi-hop questions that require systematic reasoning over entity combinations unseen during training. that require systematic reasoning over entity combinations unseen during training. Knowledge graph question answering is precisely this kind of task: the correct answer to a novel three-hop question requires chaining entities and relations in a configuration that almost certainly did not appear as a single training example. Mallen *et al.* (2023) corroborate this by showing that even GPT-4 hallucinates at a non-trivial rate on questions involving tail entities with low training-set frequency, regardless of the model's overall capability.

These findings collectively demonstrate that hallucination in large language models is not a limitation of scale, but a fundamental structural flaw. Autoregressive generation inherently lacks a mechanism to ground generated tokens in verifiable external sources, an absence that cannot be resolved simply by increasing parameter counts. The following sections review approaches that attempt to address this limitation, progressing from soft grounding via retrieval augmentation to the hard constraint mechanisms that motivate this dissertation.

2.3 Chain-of-Thought Reasoning, Self-Consistency, and the Faithfulness Ceiling of Unconstrained Multi-Step Generation

A natural question often arises when evaluating large language models: if chain-of-thought prompting already yields such significant improvements in multi-step reasoning, why do we need constrained decoding at all? The answer developed in this section is that while chain-of-thought and its extensions improve reasoning *accuracy in expectation*, they cannot offer a *structural faithfulness guarantee*, the exact property that DCA-Trie is designed to provide. Recognizing this distinction is essential, as it highlights why prompting alone is insufficient and why constrained decoding remains fundamentally necessary.

2.3.1 Chain-of-Thought Prompting: Eliciting Step-by-Step Reasoning from LLMs

Chain-of-thought (CoT) prompting, introduced by Wei *et al.* (2022), is the dominant paradigm for improving multi-step reasoning in large language models. The mechanism is straightforward: by including worked examples in the prompt that show intermediate reasoning steps leading to the final answer, the model learns to produce its own intermediate steps before

giving a final answer to a new question. This shifts the model from direct answer prediction to the decomposition of complex problems into a chain of simpler sub-problems, each resolved in sequence. Wei *et al.* (2022) show that this ability is emergent: it is negligible in models below approximately 100 billion parameters and arises sharply beyond this threshold. Empirically, CoT prompting yields substantial improvements on mathematical reasoning benchmarks such as GSM8K and MATH, as well as on multi-hop question answering tasks, in comparison with direct prompting baselines.

Formally, a CoT output for a given question x is structured as a sequence of intermediate reasoning steps followed by a final answer:

$$\big(r_1,\, r_2,\, \ldots,\, r_k,\, a\big) \tag{2.2}$$

where r_i denotes each intermediate reasoning step and a is the final answer. The language model generates each step autoregressively, using the same probabilistic mechanism described in Equation (2.1):

$$P(r_i \mid r_{<i},\, x)$$

The intrinsic limitation of this approach is immediately apparent: there is no mechanism to ensure that any given step r_i corresponds to an actual, verified triplet within the knowledge graph $\mathcal{G}$. If the model hallucinates an intermediate step r_i, that error becomes part of the conditioning context for the next step r_{i+1}, cascading forward and irreparably compromising the final answer a. As a result, a CoT sequence can read as entirely fluent, coherent, and plausible while still containing one or more steps that lack any basis in verified facts.

2.3.2 Self-Consistency and Tree-of-Thought: Improving Reliability Without Improving Faithfulness

Two major extensions of chain-of-thought attempt to address these reliability limitations, yet neither fundamentally resolves the structural faithfulness problem. The first, self-consistency (Wang *et al.*, 2023), combats the high variability of individual CoT chains by sampling multiple independent reasoning paths and selecting the most frequent final answer (majority voting). This approach effectively marginalizes out anomalous reasoning errors and substantially boosts accuracy on benchmarks like GSM8K and multi-hop QA. However, it still falls

short of providing a hard guarantee: there is no assurance that any single sampled chain is fully grounded in verified knowledge graph facts. It is entirely possible — especially for deep, multi-hop questions where error probabilities compound — for a majority of independently generated chains to all suffer from hallucinations, resulting in an incorrect majority-vote answer.

The second major extension, Tree-of-Thought (Yao *et al.*, 2023), takes a more structured approach by framing the generation process as a search over a tree of partial reasoning steps. It uses the language model as both the generator of new steps and the evaluator that scores and prunes unpromising branches. This deliberate, planning-oriented approach significantly improves performance on complex tasks that require backtracking. Ultimately, however, Tree-of-Thought operates entirely within the model's own parametric knowledge: the node scoring and pruning decisions are made by the same system generating the steps, without ever querying an external, verified knowledge graph.

The decisive observation for both of these prompting-based methods is that they can only reduce the *statistical probability* of hallucinated reasoning; they cannot actually guarantee a zero-hallucination output. This is because neither method alters the fundamental token-level generation mechanism. At any given step, the valid token set remains the entire vocabulary $\mathcal{V}$, and every token retains a strictly positive probability of being selected. This highlights the critical distinction introduced by constrained decoding, which enforces absolute structural faithfulness by setting the probability of every invalid token to exactly zero.

Table 2.1 Comparison of chain-of-thought family methods and KG-constrained decoding.

Method	**Multi-hop Acc.**	**Faithfulness Guarantee**	**Ext. Verification**	**Parametric Reliance**
Chain-of-Thought (Wei *et al.*, 2022)	High	No	No	High
Self-Consistency (Wang *et al.*, 2023)	High	No	No	High
Tree-of-Thought (Yao *et al.*, 2023)	High	No	No	High
KG-Constrained Decoding	High	Yes	Yes	Low

2.3.3 The Faithfulness Ceiling: The Mathematical Limits of Prompting

The fundamental inadequacy of prompting arises directly from its probabilistic formulation. All prompting strategies, whether relying on few-shot examples, chain-of-thought decomposition, self-consistency aggregation, or tree-structured search, operate strictly by modifying the

input provided to the language model. While this favorably shifts the probability distribution toward correct sequences, it leaves the underlying autoregressive generation mechanism entirely untouched. Because decoding still occurs over the unrestricted vocabulary $\mathcal{V}$, every token technically retains a strictly positive probability of being sampled at any given step.

True structural faithfulness, defined as the absolute guarantee that every generated token forms part of a verified knowledge graph path, requires that all invalid tokens be assigned a probability of *exactly* zero. This cannot be accomplished at the prompt level; it necessitates a fundamental intervention in the decoding algorithm itself, which is precisely the function of constrained decoding (Geng *et al.*, 2023; Willard and Louf, 2023).

As formally distinguished by Geng *et al.* (2023) and Willard and Louf (2023), prompting acts as a *soft influence*, whereas constrained decoding operates as a *hard mechanism*. While intelligent prompting can significantly elevate factual accuracy in expectation, only constrained decoding can guarantee structural faithfulness with a probability of one.

2.4 Knowledge Graph Question Answering: Structured Reasoning, Benchmark Design, and the Exponential Complexity of Multi-Hop Inference

This section establishes the technical domain in which the DCA-Trie framework operates. The principles and structural constraints introduced here provide the foundational rationale for the architectural design decisions detailed in Chapter 3.

2.4.1 Knowledge Graphs as Structured Repositories of Verified Facts: Formal Definition and Properties

A knowledge graph is formally defined as a directed labelled graph $\mathcal{G} = (\mathcal{E}, \mathcal{R}, \mathcal{T})$, where $\mathcal{E}$ is the set of entities, $\mathcal{R}$ is the set of relation types, and $\mathcal{T} \subseteq \mathcal{E} \times \mathcal{R} \times \mathcal{E}$ is the set of verified factual triplets. A standard concrete example from Freebase (Bollacker *et al.*, 2008) is the triplet (*Marie Curie*, *award_received*, *Nobel_Prize_in_Physics*), which logically connects two entities via a directed relation.

Freebase encodes tens of millions of such triplets spanning geography, history, entertainment, and science. It serves as the underlying factual resource for both benchmark datasets evaluated in this work: WebQSP (Yih *et al.*, 2016) and ComplexWebQuestions (Talmor and

Berant, 2018). Given an initial entity e_0, a multi-hop reasoning path of length L within $\mathcal{G}$ is represented as the sequence:

$$p = (e_0,\, r_1,\, e_1,\, r_2,\, e_2,\, \ldots,\, r_L,\, e_L) \tag{2.3}$$

where each consecutive step $(e_{i-1},\, r_i,\, e_i)$ forms a valid triplet within $\mathcal{T}$. In the context of KGQA, the correct answer to a multi-hop question q corresponds to the terminal entity e_L of a specific, structurally sound path p. This formalization is load-bearing throughout this dissertation: the structural faithfulness guarantee encoded by DCA-Trie strictly mandates that every generated sequence inherently forms a valid prefix of some defined path p within $\mathcal{G}$.

The critical distinction between a structured KG and an unstructured text corpus is that every triplet in $\mathcal{T}$ is human-verified, machine-readable, and unambiguously delineated. This innate topology enables KG-constrained generation to operate in a manner entirely distinct from standard retrieval augmentation (RAG). Rather than relying on a probabilistic semantic similarity search against raw text, a constraint oracle can systematically verify any candidate token against $\mathcal{G}$ via precise, deterministic graph lookups.

Concrete Example: A Knowledge Graph Subgraph from Freebase

To make the formal definition concrete, consider a small subgraph extracted from Freebase. Figure 2.2 illustrates a directed graph with six entity nodes connected by five labelled relation edges. Each edge represents a single verified triplet $(e, r, e') \in \mathcal{T}$. For instance, the triplet (*Blue_Hawaii*, *film.film.featured_film_location*, *Hawaii*) encodes the fact that the film "Blue Hawaii" was filmed in Hawaii. A valid multi-hop reasoning path through this subgraph might be:

(*Blue_Hawaii*, *film.film.featured_film_location*, *Hawaii*, *location.location.adjoining_water*, *Pacific_Ocea*

This path has length $L = 2$ (two hops) and terminates at the entity *Pacific_Ocean*. If a question asks "What large body of water adjoins the location where Blue Hawaii was filmed?", the correct answer corresponds to the terminal entity of this verified path. The structural faithfulness guarantee of DCA-Trie ensures that every intermediate step of the

model's generated reasoning chain corresponds to an actual verified triplet in $\mathcal{T}$ making hallucinated reasoning steps mathematically impossible.

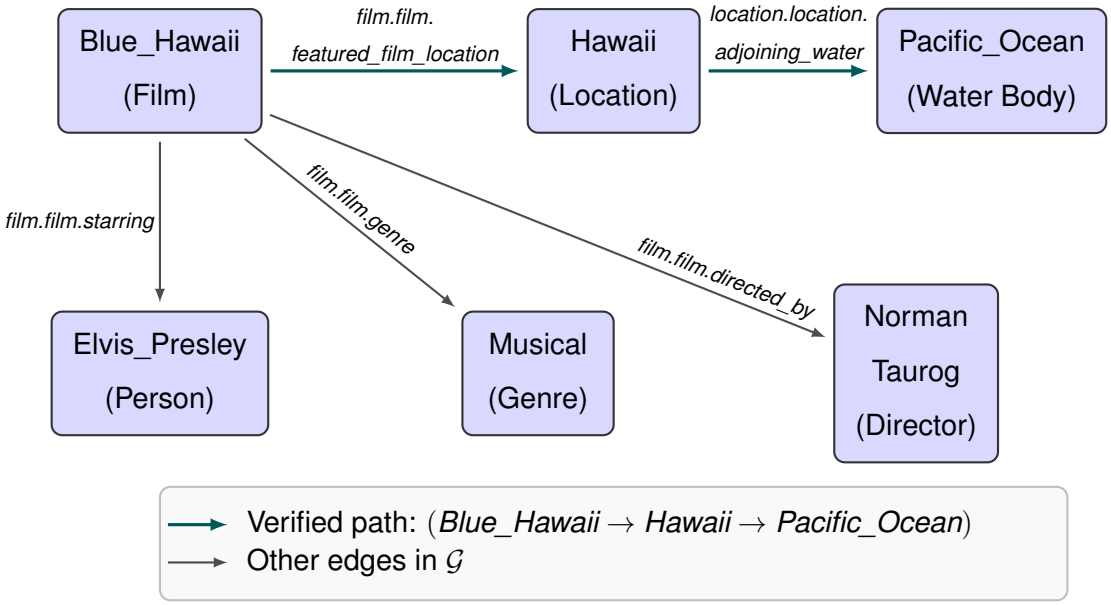

Figure 2.2 Example knowledge graph subgraph from Freebase

2.4.2 WebQSP and ComplexWebQuestions: Benchmark Architecture, Evaluation Protocol, and Hop-Depth Stratification

Two benchmark datasets provide the empirical evaluation framework for this dissertation. WebQSP (Yih *et al.*, 2016) is derived from the WebQuestions dataset and comprises 4,737 questions grounded in Freebase, with reasoning depths of one to two hops. Each question is annotated with its corresponding SPARQL query against Freebase and is typically evaluated using Hits@1 and F1 scores. Within the context of constrained generation, WebQSP represents a low-complexity verification regime; because the valid path set branching from any question entity is strictly bounded, the target path is relatively straightforward to discriminate from incorrect alternatives.

ComplexWebQuestions (CWQ; Talmor and Berant 2018) extends this paradigm by applying systematic transformations – specifically composition, conjunction, and comparative or superlative operations – to WebQSP seed queries. This yields approximately 34,000 questions requiring deeper reasoning chains spanning up to four hops. For generative frameworks, CWQ represents the combinatorial frontier where the *permissiveness problem* is most acute. At a depth of four hops, the set of structurally valid paths reachable from the question entities grows exponentially. Forced to discriminate among thousands of formally valid but

semantically irrelevant paths relying solely on its internal parametric knowledge, the LLM is exposed to the exact hallucination failure mode that DCA-Trie targets. To systematically measure this degradation, the hop-depth-stratified evaluation protocol from Lan *et al.* (2022) is adopted for the experiments in Chapter 4.

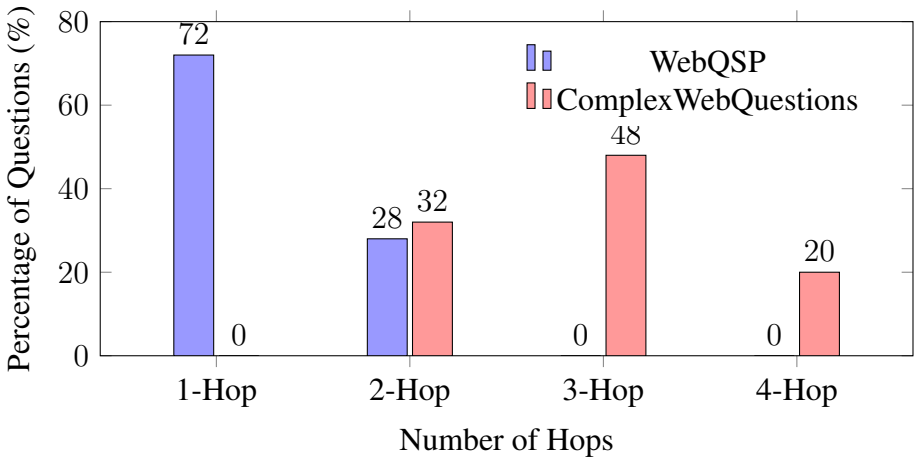

Figure 2.3 Hop-depth distributions for WebQSP and ComplexWebQuestions.

2.4.3 Multi-Hop Reasoning Complexity: Exponential Path Growth and the Permissiveness Problem Quantified

The qualitative difference between one-hop and multi-hop questions can be made precise through a combinatorial analysis. A one-hop question requires selecting a single correct relation from the set of outgoing edges of the question entity e_q. The valid candidate set is bounded by the out-degree d of e_q. A multi-hop question of depth L requires selecting a correct sequence of L consecutive steps, where at each step the valid set multiplies. The upper bound on the number of valid KG paths of length L reachable from the question entities $\mathcal{E}_q$ is:

$$|\mathcal{P}_L| \;\leq\; |\mathcal{E}_q| \times d^L \tag{2.4}$$

where d represents the average out-degree of entities in $\mathcal{G}$. Assuming a conservative average out-degree $d \approx 20$ for well-connected entities in Freebase, and a typical query linking to $|\mathcal{E}_q| \in \{1, 2, 3\}$ seed entities, we can instantiate Equation (2.4): at $L = 2$, the upper bound is $3 \times 400 = 1{,}200$ valid paths; at $L = 3$, the bound grows to $3 \times 8{,}000 = 24{,}000$ valid paths;

at $L = 4$, the bound reaches $3 \times 160{,}000 = 480{,}000$ valid paths. In each case, exactly one path is the correct targeted answer.

This exponential growth formally quantifies the permissiveness problem inherent to GCR and DoG. When the constraint oracle relies purely on structural validity, $f(\mathcal{G}, \mathcal{E}_q)$, the valid token set at every generation step is flooded with thousands of structurally permitted but semantically irrelevant paths. The LLM is consequently forced to discriminate among them relying solely on its internal parametric knowledge – precisely the hallucination-prone failure mode that constrained decoding was originally designed to suppress (Luo *et al.*, 2024; Li *et al.*, 2024). The Semantic Irrelevance Ratio (SIR), introduced in Chapter 3 of this dissertation, provides the first direct measure of this phenomenon. Empirical evidence for performance degradation as a function of hop depth is documented in He *et al.* (2021b) and Lan *et al.* (2022).

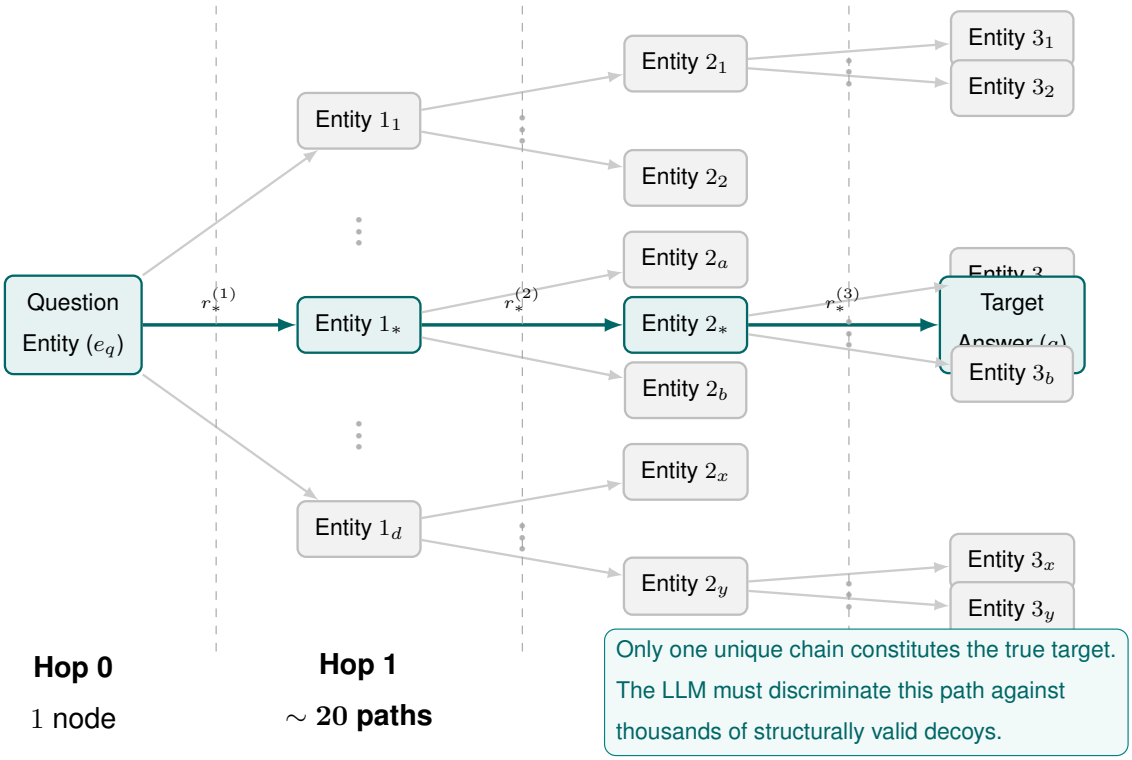

Figure 2.4 Exponential growth of structurally valid KG paths with hop depth.

2.5 Constrained Decoding for Faithful Text Generation: From Logit Masking to Knowledge Graph Oracles

This section establishes the technical lineage of constrained decoding paradigms leading to the proposed framework. It progresses from the foundational logit-masking mechanism through grammar-based constraints to contemporary knowledge-graph-specific oracles. Each framework reviewed serves either as a direct methodological baseline or as a foundational model whose specific structural limitations the DCA-Trie architecture is designed to resolve.

2.5.1 The Logit-Masking Paradigm: Hard Constraints on Autoregressive Vocabulary Selection

The logit-masking paradigm is the foundational mechanism for hard-constrained generation (Geng *et al.*, 2023; Willard and Louf, 2023). The key insight is that rather than influencing the probability distribution over tokens, as prompting does, a constraint oracle computes a binary mask that sets the logit of every invalid token to negative infinity (effectively zero probability after softmax) before the token selection step. This renders the generation of invalid tokens mathematically impossible rather than merely improbable.

Formally, let $\boldsymbol{\alpha}_t \in \mathbb{R}^{|\mathcal{V}|}$ be the raw logit vector at decoding step t, and let $\mathbf{m}_t \in \{0,1\}^{|\mathcal{V}|}$ be the binary validity mask computed by the constraint oracle. The masked logit vector is:

$$\tilde{\boldsymbol{\alpha}}_t \;=\; \mathbf{m}_t \odot \boldsymbol{\alpha}_t \tag{2.5}$$

where $\odot$ denotes element-wise product (with the convention that $0 \odot \alpha = -\infty$ before softmax). Token selection under greedy decoding is $y_t = \arg\max(\tilde{\boldsymbol{\alpha}}_t)$. Under beam search with beam width k, the top-k partial sequences are maintained and ranked by cumulative log-probability of masked logits.

The central research question of this dissertation follows directly from this formulation: *how should* $\mathbf{m}_t$ *be computed?* The answer determines the expressiveness, faithfulness, and efficiency of the entire framework. GCR computes $\mathbf{m}_t$ from a function $f(\mathcal{G}, \mathcal{E}_q)$ evaluated once before generation. DCA-Trie computes $\mathbf{m}_t$ from a function $f(\mathcal{G}, \mathcal{E}_q, q, y_{<t})$ updated at every step. Every subsequent subsection is a step in the progression toward this answer.

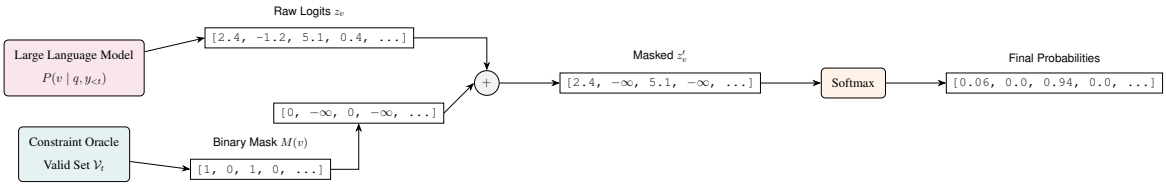

Figure 2.5 The logit-masking mechanism. Token selection operates exclusively over the valid, non-zero subset defined by the constraint oracle mask.

2.5.2 Grammar-Constrained Decoding: GCD and Outlines as Foundational Format Constraints

Geng *et al.* (2023) introduce Grammar-Constrained Decoding (GCD), which constrains generation to strings conforming to a context-free grammar by using an Earley parser as the constraint oracle. Willard and Louf (2023) introduce Outlines, which achieves the same goal for regular-expression and JSON-schema constraints via finite-state automata. Both works establish the critical efficiency insight that the constraint structure can be precomputed offline, as parse states or automaton transitions, and queried in amortized constant time at each inference step. This precomputation insight is central to DCA-Trie v1, which preserves it through static semantic filtering at trie construction time.

Algorithm 1 Grammar-Constrained Decoding Algorithm

Require: Context-free grammar $\mathcal{G}_{\text{CFG}}$; LLM with vocabulary $\mathcal{V}$; prompt x

Ensure: Output sequence $y_1 \dots y_T$ valid under $\mathcal{G}_{\text{CFG}}$

1: **Precompute:** Earley parse automaton $\mathcal{A}$ from $\mathcal{G}_{\text{CFG}}$
2: **for** $t = 1$ **to** T **do**
3: $\quad \mathcal{V}_t \leftarrow \{v \in \mathcal{V} : y_{<t} \| v \text{ is a valid grammar prefix under } \mathcal{A}\}$ ▷ Oracle query
4: $\quad \mathbf{m}_t[v] \leftarrow \mathbf{1}[v \in \mathcal{V}_t]$ for all $v \in \mathcal{V}$
5: $\quad \tilde{\boldsymbol{\alpha}}_t \leftarrow \mathbf{m}_t \odot \textsc{LLMLogits}(y_{<t}, x)$
6: $\quad y_t \leftarrow \arg\max(\tilde{\boldsymbol{\alpha}}_t)$
7: **end for**
8: **return** $y_1 \dots y_T$

Figure 2.6 Grammar-Constrained Decoding Algorithm.

The shared limitation of GCD and Outlines is stated precisely: both frameworks constrain the *format* of the output (grammar, schema, structure) but not the *factual content*. A grammatically valid JSON response can contain entirely hallucinated facts. As annotated in Figure 2.6, the valid set $\mathcal{V}_t$ depends only on the partial string and the grammar; there is no dependence on question semantics or on a verified factual source. This is the same structural limitation present in GCR's KG-Trie, which constrains to KG path format while remaining semantically blind to the question. The transition from format constraints to knowledge-graph content constraints is the next step in the progression.

2.5.3 Graph-Constrained Reasoning (GCR): Structural Faithfulness via KG-Trie Oracles

Graph-Constrained Reasoning (GCR; Luo *et al.* 2024) is the primary baseline for this dissertation and represents the state of the art in KG-constrained decoding at the time of this work. GCR introduces five components. First, *entity linking* identifies the question entity set $\mathcal{E}_q$ from the natural language question before generation begins. Second, *KG-Trie construction* performs a breadth-first search from $\mathcal{E}_q$ up to L hops, encoding all reachable KG paths as a prefix tree. A trie node at depth d corresponds to a partial path of d tokens. Third, *constrained decoding* queries the trie at each generation step with the current partial generation $y_{<t}$ to return the set of tokens that extend a valid trie prefix. Fourth, the *faithfulness guarantee* follows: because every token must correspond to a valid trie prefix, and every trie path corresponds to a verified triplet chain in $\mathcal{T}$, structural hallucination is mathematically impossible. Fifth, GCR achieves 92.6 Hits@1 on WebQSP and 75.8 Hits@1 on CWQ with verified 100% structural faithfulness using a Llama-3.1-8B backbone (Luo *et al.*, 2024).

Formally, the valid token set under GCR is:

$$\mathcal{W}_{\text{val}}^{\text{GCR}} \;=\; f(\mathcal{G},\, \mathcal{E}_q) \tag{2.6}$$

The KG-Trie $\mathcal{T}_{\text{trie}}$ is constructed as $\mathcal{T}_{\text{trie}} = \text{BFS}(\mathcal{G}, \mathcal{E}_q, L)$. At each step t:

$$\mathcal{V}_t = \left\{ v \in \mathcal{V} : \text{prefix}(y_{<t}) \| v \in \mathcal{T}_{\text{trie}} \right\}$$

The critical annotation is that $\mathcal{W}_{\text{val}}^{\text{GCR}}$ has *no dependence* on q (the question text) or on $y_{<t}$ (the partial generation state). The valid set is semantically blind and computationally

fixed before generation begins. This is the exact formulation that DCA-Trie upgrades to $\mathcal{W}_{\text{val}}^{\text{DCA}}(t) = f(\mathcal{G}, \mathcal{E}_q, q, y_{<t})$.

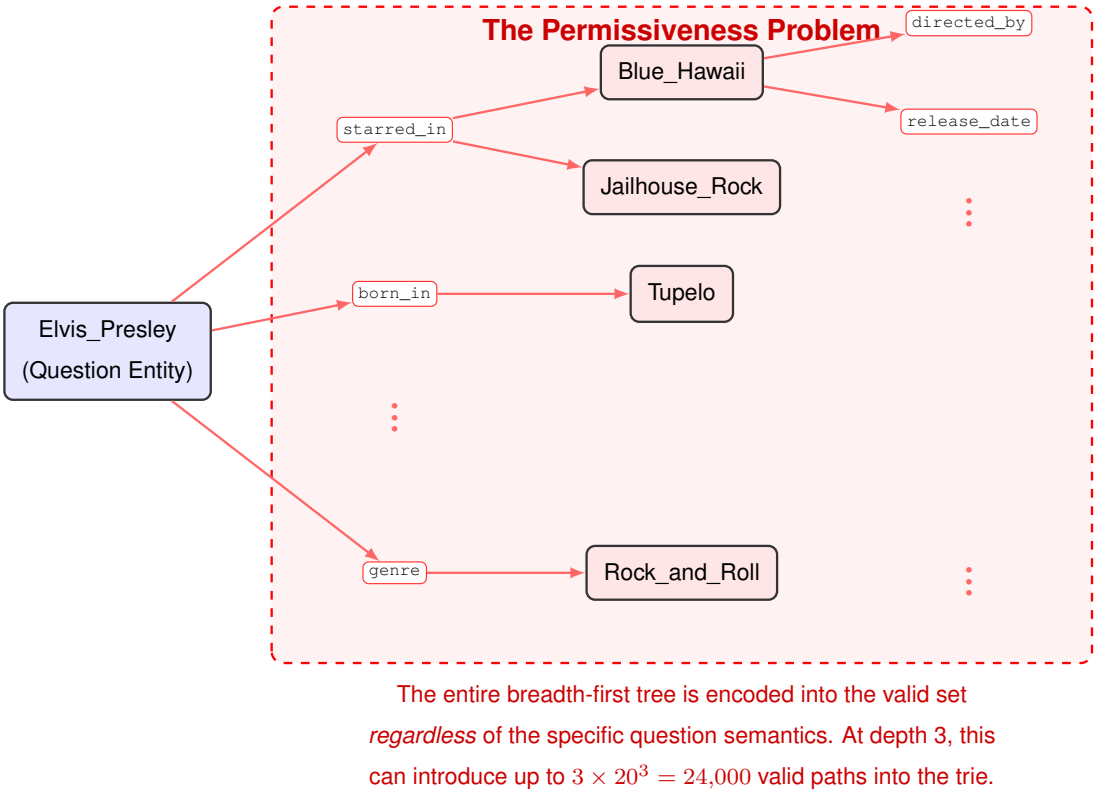

Figure 2.7 The GCR KG-Trie illustrating the permissiveness problem.

2.5.4 Decoding on Graphs (DoG): Step-Wise Topological Expansion and Its Persistent Limitations

Li *et al.* (2024) introduce Decoding on Graphs (DoG) as a partial remedy to GCR's static trie. DoG introduces step-wise trie expansion: after committing to an entity e_t during generation, the valid set is extended with all KG neighbors of e_t. Formally, after committing entity e_t:

$$\mathcal{T}_{t+1} \;=\; \mathcal{T}_t \;\cup\; \big\{(e_t,\, r,\, e') : (e_t,\, r,\, e') \in \mathcal{T} \text{ for all } r, e'\big\} \tag{2.7}$$

This makes the constraint topologically dynamic: it grows as reasoning progresses through the graph. DoG offers two genuine improvements over GCR. First, the valid set at step t reflects the entities already committed to in the reasoning chain, so the trie cannot contain paths that are structurally inconsistent with the committed reasoning direction. Second, the

dynamic expansion can handle knowledge graphs too large to enumerate statically within the pre-generation BFS.

However, DoG has two persistent limitations. First, as is explicit in Equation (2.7), the expansion rule is semantically blind: every neighbor of every visited entity is admitted to $\mathcal{T}_{t+1}$ regardless of its relevance to q or to the partial generation $y_{<t}$. The permissiveness problem therefore persists: the number of paths in the valid set may remain in the thousands at every step. Second, DoG sacrifices the computational efficiency of precomputed trie structures identified by Willard and Louf (2023): because the trie must be rebuilt at each step, the amortized per-step query cost of precomputed structures is forfeited.

DCA-Trie addresses both limitations simultaneously. The DCA-Trie v2 expansion rule replaces the semantically blind expansion of Equation (2.7) with a semantically conditioned rule (defined formally in Chapter 3):

$$\mathcal{T}_{t+1} = \mathcal{T}_t \cup \left\{ (e_t,\, r,\, e') : \textsc{score}(e',\, q,\, y_{<t}) \geq \tau \right\}$$

The single structural change, conditioning expansion on a semantic relevance score, is the core algorithmic contribution of DCA-Trie v2.

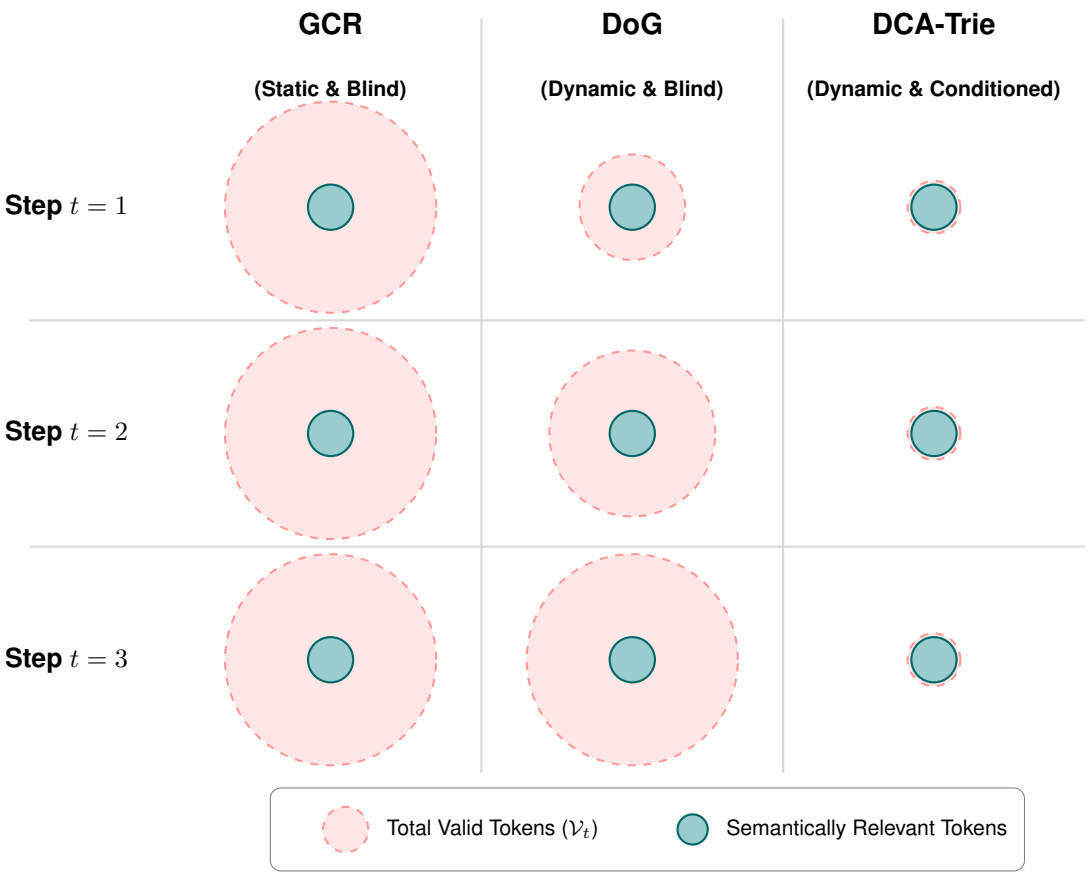

Figure 2.8 Evolution of the constraint valid set across decoding steps.

2.6 Beam Search Decoding and Prefix-Tree Constraint Integration: Mechanics, Efficiency, and the Step-Wise Constraint Problem

DCA-Trie v2 operates step-wise during beam search. Reviewers evaluating the method chapter need to understand both the beam search mechanism and how prefix-tree constraints interact with beam expansion before the DCA-Trie design can be followed. This section provides that foundation.

2.6.1 Beam Search: Approximate Global Optimization Over Autoregressive Sequences

Beam search is the standard approximate decoding algorithm for autoregressive sequence generation (Holtzman *et al.*, 2020). The algorithm maintains a *beam* (a set of k partial sequences) at each generation step. At step t, every beam hypothesis is expanded by every vocabulary token, masked logit scores are computed, and only the top-k hypotheses by cumulative log-probability are retained. Formally, for each beam hypothesis $h \in \mathcal{B}_t$:

$$\text{SCORE}(h\|v) = \text{SCORE}(h) + \log \text{SOFTMAX}\big(\tilde{\boldsymbol{\alpha}}_t\big)[v], \quad \forall\, v \in \mathcal{V} \text{ with } \tilde{\boldsymbol{\alpha}}_t[v] > 0 \tag{2.8}$$

The updated beam is $\mathcal{B}_{t+1} =$ top-k over all (h, v) pairs by $\text{SCORE}(h\|v)$. The beam width k controls the exploration-exploitation tradeoff: $k = 1$ reduces to greedy decoding; larger k explores a wider portion of the generation space at the cost of k-fold compute per step.

When combined with the logit mask from Equation (2.5), beam search operates exclusively within the valid subspace defined by the constraint oracle. The KG-Trie defines this subspace structurally; DCA-Trie's semantic filter narrows it further, reducing the number of (h, v) candidate pairs that must be scored at each beam expansion step. This reduction is what makes DCA-Trie computationally competitive with GCR despite the additional scoring overhead.

2.6.2 Prefix Trees as Constraint Oracles: Efficiency, Precomputation, and the Dynamic Rebuilding Cost

A prefix tree (trie; Fredkin 1960) over KG paths supports $O(d)$ lookup at each generation step: given the current partial sequence $y_{<t}$, the set of valid next tokens is exactly the set of children of the current trie node. This is the efficiency insight of Willard and Louf (2023): precomputing the trie amortizes the constraint lookup cost across all generation steps, making constrained decoding practical at inference time. The trie construction complexity is $O(|\mathcal{P}| \times L)$, where $|\mathcal{P}|$ is the number of valid paths encoded and L is the maximum path length.

DoG's step-wise expansion sacrifices this benefit by rebuilding the trie at each step at a cost of $O(|\mathcal{N}_t| \times d)$ per step, where $|\mathcal{N}_t|$ is the number of newly reachable entities at step t. While not prohibitive per step, this cost accumulates over L steps and eliminates the precomputation advantage entirely. DCA-Trie v1 (static semantic filtering at construction time) preserves the precomputation benefit entirely, since the filtered trie is built once before generation. DCA-Trie v2 (dynamic step-wise expansion with semantic filtering) incurs per-step rebuilding overhead similar to DoG, but the semantic filter prunes most candidate expansions, so the effective expansion set $|\mathcal{N}_t^{\text{filtered}}|$ is substantially smaller than $|\mathcal{N}_t|$ in DoG.

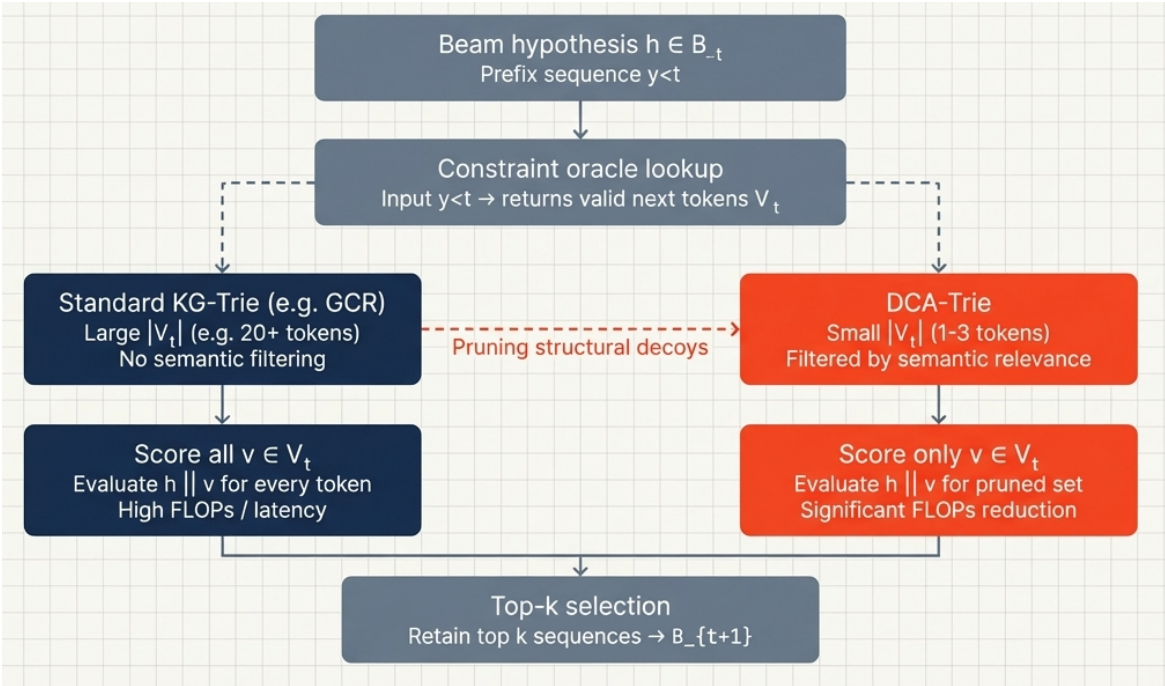

Figure 2.9 Interaction between prefix-tree constraint lookup and beam search expansion.

Image by Authors

2.6.3 The Step-Wise Constraint Problem: When the Constraint Must Know What Has Been Generated

The step-wise constraint problem is the core design challenge motivating DCA-Trie. In GCR, the constraint oracle is computed once before generation and queried at each step without memory of what has been generated beyond the structural trie path. In DCA-Trie v2, the oracle is updated after each committed entity to reflect the semantic direction of the current reasoning chain. This requires the oracle to maintain state across generation steps; specifically, the partial generation $y_{<t}$ and the question embedding must be re-incorporated at each update.

The formal requirement is:

$$\mathcal{W}_{\text{val}}^{\text{DCA}}(t) \;=\; f\big(\mathcal{G},\, \mathcal{E}_q,\, q,\, y_{<t}\big) \tag{2.9}$$

where the function f must satisfy three conditions. (i) *Structural faithfulness*: every token $v \in \mathcal{W}_{\text{val}}^{\text{DCA}}(t)$ corresponds to a valid KG path prefix in $\mathcal{G}$. (ii) *Semantic relevance*: the proportion of $\mathcal{W}_{\text{val}}^{\text{DCA}}(t)$ consisting of semantically irrelevant paths is strictly less than that

of $\mathcal{W}_{\text{val}}^{\text{GCR}}$. (iii) *Recall preservation*: the gold answer path is never excluded from $\mathcal{W}_{\text{val}}^{\text{DCA}}(t)$, with false negative rate held below 5% on the WebQSP validation set. These three conditions constitute the formal specification of DCA-Trie, evaluated quantitatively in Chapter 4.

2.7 Retrieval-Augmented Generation: Architectural Variants, KG-Augmented Approaches, and the Hard Limit of Soft Grounding

Retrieval-augmented generation is the dominant paradigm for grounding LLM outputs in external knowledge. This section reviews the foundational RAG architecture, KG-specific retrieval variants, and the latest advanced RAG systems, before providing a formal impossibility argument establishing that no RAG variant can achieve the structural faithfulness guarantees that constrained decoding provides.

2.7.1 The RAG Paradigm: Dense Retrieval, Fusion-in-Decoder, and the Soft Grounding Ceiling

Lewis *et al.* (2020) introduce Retrieval-Augmented Generation, in which a dense retriever selects relevant passages from an external knowledge source based on embedding similarity, and a reader LLM generates the answer conditioned on both the question and the retrieved passages. The retriever is typically implemented as a dual-encoder (Dense Passage Retrieval; Karpukhin *et al.* 2020), computing embedding similarity between the question and candidate passages offline. Fusion-in-Decoder (Izacard and Grave, 2021) extends the original RAG architecture by encoding all retrieved passages independently through the encoder and fusing their representations in the decoder, improving coverage across a larger set of retrieved documents.

The critical shared limitation of these approaches is architectural: retrieved context is a *soft influence*: the LLM is free to generate any token regardless of the retrieved material. The retriever controls what the LLM *sees*; the constraint oracle controls what the LLM may *generate*. RAG addresses the former; constrained decoding addresses the latter. Agrawal *et al.* (2024a) document the concrete consequence: even with explicit KG access, RAG-augmented models produce hallucinated reasoning steps in a substantial fraction of cases.

2.7.2 KG-Augmented Retrieval: Structured Context Without Structural Guarantees

Several works augment the RAG paradigm with KG subgraphs rather than text passages, bringing the approach closer in spirit to constrained KG decoding. He *et al.* (2021b) use attention over retrieved KG subgraphs to guide multi-hop reasoning chains. Yasunaga *et al.* (2021) combine text and KG context via graph neural networks, allowing the model to attend jointly to retrieved passages and KG neighbourhoods. Shi *et al.* (2021) retrieve KG paths as natural language context strings.

For each of these approaches, the faithfulness limitation is identical: the retrieved KG subgraph or path is provided to the LLM as additional input context. The LLM may attend to it, partially ignore it, or generate claims that go beyond it. The KG is advisory, not constraining. This is the decisive architectural distinction from GCR and DCA-Trie, where the KG is *mechanically* enforced through the logit mask rather than provided as contextual suggestion.

2.7.3 Advanced RAG Variants: Self-RAG, FLARE, IRCoT, and GraphRAG

Four advanced RAG variants warrant individual analysis as potential alternatives to constrained decoding that a reviewer may raise.

Self-RAG (Asai *et al.*, 2024) augments the standard RAG pipeline by inserting special reflection tokens that allow the model to decide dynamically when retrieval is needed and whether retrieved passages are relevant or supportive. This improves retrieval precision and citation accuracy. However, Self-RAG does not modify the token selection mechanism: the LLM remains free to generate any token at any step, and the reflection tokens are themselves generated autoregressively. Unlike DCA-Trie, Self-RAG cannot guarantee structural faithfulness.

FLARE (Jiang *et al.*, 2023) uses the LLM's next-token prediction probability as a signal to trigger retrieval when confidence falls below a threshold. This addresses the question of *when* to retrieve but leaves the generation faithfulness problem untouched: retrieved context is again a soft input influence.

IRCoT (Trivedi *et al.*, 2023) interleaves retrieval with chain-of-thought reasoning steps, performing a retrieval call at each CoT step. This substantially improves multi-hop accuracy

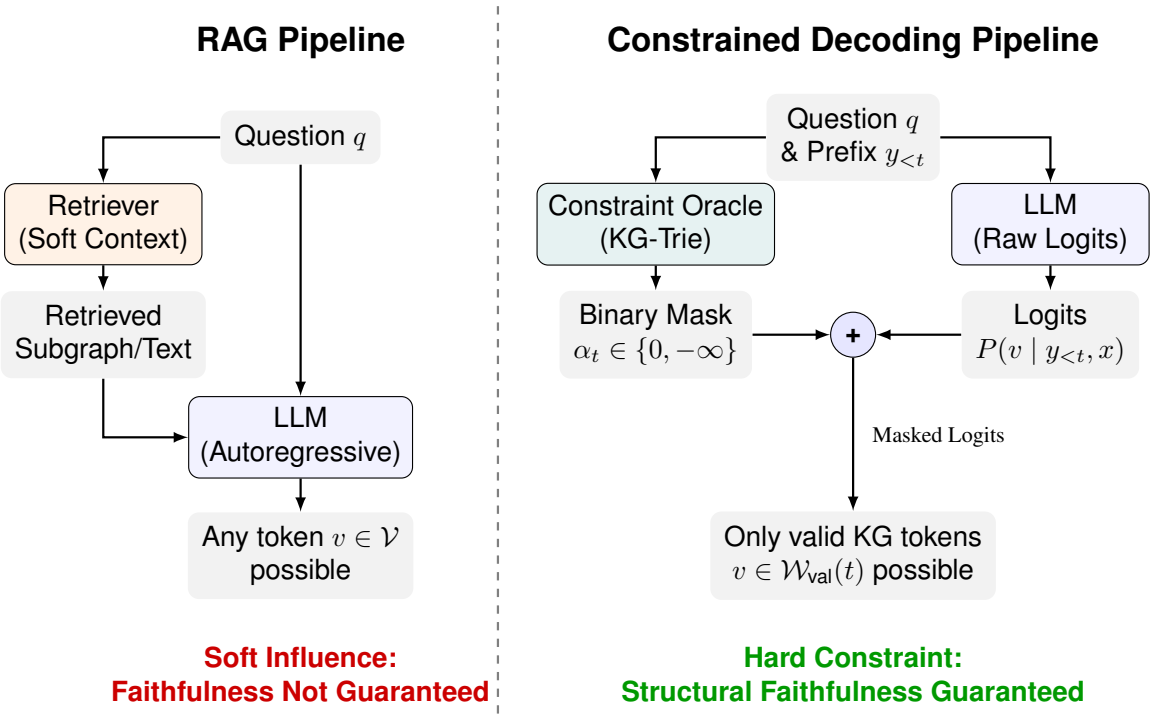

Figure 2.10 Architectural comparison of standard Retrieval-Augmented Generation (left) and KG-constrained decoding (right).

by ensuring that each reasoning step is informed by freshly retrieved context. Each CoT step nevertheless remains an unconstrained autoregressive generation, retaining strictly positive probability over all vocabulary tokens.

GraphRAG (Edge *et al.*, 2024) is the strongest RAG-based competitor to DCA-Trie, as it explicitly retrieves community-structured graph context to support complex relational queries. GraphRAG retrieves graph context but does not constrain decoding: a generation conditioned on GraphRAG context can still produce tokens that correspond to no KG triplet in $\mathcal{T}$.

For every variant, the same single sentence applies: unlike DCA-Trie, none of these methods modifies the token selection mechanism and therefore none can guarantee structural faithfulness.

2.7.4 The Formal Impossibility of Structural Faithfulness Under Soft Grounding

The following argument establishes that no retrieval-augmented approach can provide structural faithfulness guarantees as defined in this dissertation.

Claim: Any approach that grounds LLM generation by modifying its input (retrieved context, subgraphs, passages) without modifying the logit computation mechanism cannot guarantee

structural faithfulness with probability one.

Proof sketch: Structural faithfulness requires that every generated token y_t corresponds to a verified KG triplet. This requires that every token $v \notin \mathcal{W}_{\text{val}}(t)$ has exactly zero probability at step t, i.e., $P(v \mid y_{<t}, x) = 0$ for all invalid v. Under any RAG approach, the LLM's raw logit distribution assigns strictly positive logit to all tokens in $\mathcal{V}$ (or at least to all tokens not explicitly set to $-\infty$). Retrieval augmentation modifies the input x to include retrieved context $\hat{x}$, but does not modify the logit computation $P(\cdot \mid y_{<t}, \hat{x})$: the distribution remains over the full vocabulary with strictly positive probabilities. Therefore, for any RAG approach, there exists strictly positive probability of generating a non-KG-valid token at any step. Probability-one structural faithfulness is not achievable under this architecture. The only mechanism that achieves it is a constraint oracle that explicitly sets $\alpha_t[v] = -\infty$ for all invalid v: the logit-masking mechanism of Equation (2.5), as implemented in GCR, DoG, and DCA-Trie. Luo *et al.* (2024) provide empirical confirmation: GCR achieves 100% structural faithfulness where RAG-based approaches do not. □

2.8 Semantic Relevance Scoring with Sentence Transformer Embeddings: Training Objectives, Architecture Selection, and Application to Constraint Oracle Construction

This section justifies every design choice in DCA-Trie's semantic relevance scorer, responding in advance to the reviewer objections: "Why not BM25?", "Why not use the LLM's own hidden states?", and "Why all-MiniLM-L6-v2 over a stronger encoder?"

2.8.1 Sentence-BERT: Training Objective, Architecture, and the Semantic Similarity Property

Reimers and Gurevych (2019) introduce Sentence-BERT (SBERT), a modification of the BERT architecture trained specifically to produce dense vector representations in which semantically similar sentences have high cosine similarity. The architecture uses a BERT-like transformer encoder with mean-pooling over token representations to produce a fixed-length sentence embedding. SBERT is trained using a Siamese network on labelled sentence pairs. The classification-objective training loss for a pair (s_1, s_2) with embeddings $\mathbf{u} = E(s_1)$ and $\mathbf{v} = E(s_2)$ is:

$$\mathcal{L} = \textsc{CrossEntropy}\Big(\text{softmax}\big(\mathbf{W}_t \cdot (\mathbf{u},\ \mathbf{v},\ |\mathbf{u} - \mathbf{v}|)\big),\ y\Big) \tag{2.10}$$

where $(\mathbf{u}, \mathbf{v}, |\mathbf{u}-\mathbf{v}|)$ is the concatenation of the two embeddings with their element-wise absolute difference, $\mathbf{W}_t$ is a trainable classification head, and $y \in \{\text{entailment}, \text{neutral}, \text{contradiction}\}$. This objective pulls embeddings of entailed pairs together and pushes contradiction pairs apart in the representation space. At inference time, $\cos(E(s_1), E(s_2))$ measures the degree of semantic alignment between two sentences without requiring the classification head, enabling cosine similarity to serve as a reliable semantic relevance proxy without any task-specific fine-tuning.

2.8.2 Encoder Architecture Selection: all-MiniLM-L6-v2 vs Alternatives

DCA-Trie requires a semantic encoder that is invoked at every beam search step during constrained decoding. The selection of the encoder architecture therefore involves a tradeoff between semantic quality and inference latency that is more stringent than standard offline retrieval settings. Table 2.2 presents a systematic comparison of five candidate architectures.

Table 2.2 Comparison of sentence encoder architectures.

Encoder	**Params**	**Emb. Dim**	**Latency (ms)**	**STS-B**	**Step-Wise?**
all-MiniLM-L6-v2 *(selected)*	22M	384	∼80	82.3	Yes
all-mpnet-base-v2	110M	768	∼320	84.6	No
BGE-M3	568M	1024	>500	87.0	No
LLM hidden states	—	varies	high	—	No
BM25 (lexical)	—	—	<5	63.0	Yes

Latency figures are approximate per-query values on a single CPU core. STS-B scores are Spearman correlations reported in the sentence-transformers library benchmarks (Reimers and Gurevych, 2019).

all-MiniLM-L6-v2 (Reimers and Gurevych, 2019) is a six-layer MiniLM model distilled from a larger teacher, producing 384-dimensional embeddings with approximately 22M parameters and roughly 80ms inference latency on CPU. It achieves strong semantic performance on STS benchmarks while maintaining the lowest inference latency among semantically competitive models.

all-mpnet-base-v2 achieves stronger semantic performance but requires approximately 110M parameters and incurs around 4× the inference latency. For iterative beam-search decoding in which the scorer is called at every step, this latency accumulates unacceptably over the full generation length.

BGE-M3 (Chen *et al.*, 2024) is a state-of-the-art multilingual retrieval encoder with stronger performance than MiniLM, but its 568M parameters and resulting inference latency make it infeasible for step-wise decoding on a single A100 GPU.

LLM hidden states as a relevance scorer would be circular: the LLM's last-layer representations encode the same parametric biases that cause hallucination. Using them as a relevance filter does not introduce any external grounding and defeats the purpose of the constraint mechanism.

BM25 is fast and interpretable but a lexical retrieval method: it cannot capture semantic similarity between paraphrase-equivalent KG paths that share the same meaning but no surface tokens, which is common in Freebase relation strings.

The conclusion is that all-MiniLM-L6-v2 provides the optimal balance of semantic quality and inference speed for iterative step-wise constraint filtering.

2.8.3 Cosine Similarity as a Hard Constraint Signal: From Soft Retrieval to Binary Admission

In standard dense retrieval, cosine similarity produces a continuous relevance ranking over candidate documents. DCA-Trie's key conceptual contribution in this component is the conversion of this continuous signal into a *binary admission criterion* that determines which tokens the LLM is permitted to generate.

Given a candidate KG path p expressed as a natural language string $(e_0\, r_1\, e_1\, r_2\, e_2 \ldots)$, the question q, and the partial generation $y_{<t}$, the semantic relevance score is:

$$\textsc{score}(p,\, q,\, y_{<t}) \;=\; \cos\!\big(E(p),\;\; E(\text{concat}(q,\, y_{<t}))\big) \tag{2.11}$$

where $E(\cdot)$ is the all-MiniLM-L6-v2 encoder. The binary admission criterion is: admit p to $\mathcal{W}_{\text{val}}^{\text{DCA}}(t)$ if and only if $\textsc{score}(p, q, y_{<t}) \geq \tau$. This threshold function converts the continuous

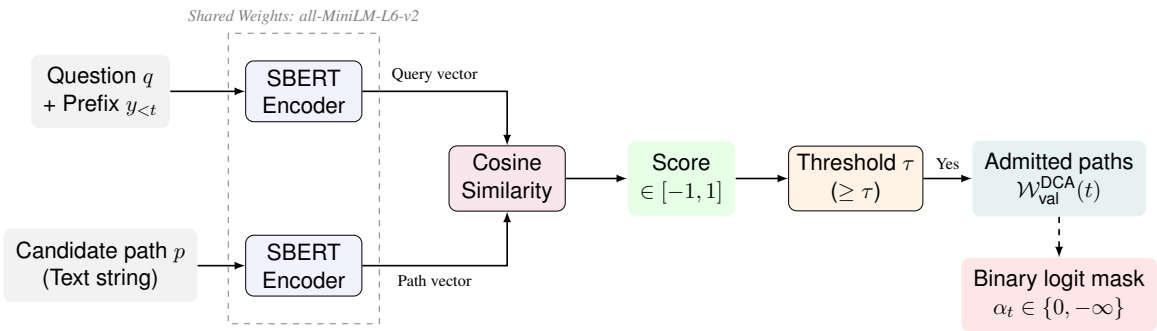

Figure 2.11 End-to-end semantic scoring and binary admission pipeline.

cosine similarity into a binary mask entry, integrating semantic relevance directly into the logit-masking framework of Equation (2.5). The threshold τ is the key hyperparameter controlling the precision-recall tradeoff of the relevance filter; its selection procedure is described in Chapter 3 and evaluated in Chapter 4.

2.9 Related Work

Sections 2.2–2.8 surveyed the landscape thematically. This section situates DCA-Trie within the most relevant prior work likely to be cited as baseline literature.

2.9.1 KG-Constrained Decoding: GCR, DoG, and the Unresolved Permissiveness Problem

GCR (Luo *et al.*, 2024) constructs a static KG-Trie from $f(\mathcal{G}, \mathcal{E}_q)$ before generation, reporting 92.6 Hits@1 on WebQSP and 100% structural faithfulness. Its valid token set, however, is semantically blind and computationally unchanged across generation steps (Section 2.5). DoG (Li *et al.*, 2024) makes the oracle dynamic, $f(\mathcal{G}, \mathcal{E}_q, t)$, by expanding the trie after each committed entity; this yields 90.99 Hits@1 on WebQSP and 76.08 on CWQ with 100% structural faithfulness, but still without semantic conditioning and without precomputation. As shown in Equation (2.7), DoG expansion remains semantically blind because it admits KG neighbors regardless of relevance.

The proposed DCA-Trie variants in Table 2.3 are positioned to address this specific gap at the token-constraint level. DCA-Trie v1 applies question-conditioned semantic filtering in a static precomputed oracle, $f(\mathcal{G}, \mathcal{E}_q, q)$. DCA-Trie v2 further adds step-wise conditioning on partial generation, $f(\mathcal{G}, \mathcal{E}_q, q, y_{<t})$, with partial precomputation. In both variants, semantic

conditioning is incorporated into the hard constraint itself while preserving the structural-faithfulness guarantee by construction; their Hits@1 values remain TBD in this chapter.

Table 2.3 Primary related work comparison.

Method	Oracle	Faithfulness	Sem. Cond.	Dynamic	Precomp.	Hits@1
GCR (Luo *et al.*, 2024)	$f(\mathcal{G}, \mathcal{E}_q)$	100%	No	No	Yes	92.6
DoG (Li *et al.*, 2024)	$f(\mathcal{G}, \mathcal{E}_q, t)$	**100%**	**No**	**Yes**	**No**	**90.99** (WebQSP) / **76.08** (CWQ)
RouterKGQA (Yuan *et al.*, 2026)	Routing + filter	Partial	Yes (answer)	Yes	N/A	+0.49 vs GCR
DCA-Trie v1 (proposed)	$f(\mathcal{G}, \mathcal{E}_q, q)$	**100%**	**Yes (token)**	**No**	**Yes**	**TBD**
DCA-Trie v2 (proposed)	$f(\mathcal{G}, \mathcal{E}_q, q, y_{<t})$	**100%**	**Yes (token)**	**Yes**	**Partial**	**TBD**

2.9.2 Hallucination Mitigation and KGQA: The Soft-to-Hard Spectrum

Soft grounding approaches constitute one cluster of related work. RAG (Lewis *et al.*, 2020), Self-RAG (Asai *et al.*, 2024), GraphRAG (Edge *et al.*, 2024), DAGQA (He *et al.*, 2021b), QA-GNN (Yasunaga *et al.*, 2021), and self-consistency (Wang *et al.*, 2023) all improve faithfulness in expectation by providing the LLM with relevant external context or by aggregating over multiple generation samples. As established formally in Section 2.7, none of these approaches modifies the token selection mechanism: they cannot guarantee structural faithfulness with probability one, because every token in $\mathcal{V}$ retains strictly positive probability at every step.

Format-constrained decoding constitutes the second cluster. GCD (Geng *et al.*, 2023), Outlines (Willard and Louf, 2023), and entity-constrained generation (De Cao *et al.*, 2021) establish the logit-masking paradigm on which DCA-Trie builds. These works constrain output format or entity vocabulary rather than multi-hop KG path faithfulness: their constraint oracle is a function of a grammar or entity list, not of a knowledge graph with relational structure. DCA-Trie inherits the logit-masking mechanism from this line of work and extends it to semantically conditioned KG path constraints.

2.9.3 Semantic Similarity for Graph Reasoning: The Retrieval-Constraint Boundary

Several works apply semantic similarity signals in graph-based reasoning. Path retrieval for KGQA (Shi *et al.*, 2021; Saxena *et al.*, 2022) uses semantic similarity to rank candidate paths for retrieval. Attention-guided graph traversal in DAGQA (He *et al.*, 2021b) uses attention weights to focus multi-hop reasoning.Semantic path reranking approaches select among candidate answers after generation. In every case, semantic similarity operates as

a *soft signal* that influences which paths the LLM sees or which answer is selected after generation; it does not apply semantic similarity as a hard binary admission criterion that determines which tokens the LLM is *permitted* to generate.

This distinction between soft signals and hard constraints defines the exact novelty boundary of DCA-Trie. No prior work combines semantic relevance scoring with hard constraint oracle construction. Works that provide structural faithfulness (GCR,DoG) do not condition on semantics. Works that use semantic similarity (path retrieval, reranking) do not provide structural faithfulness. DCA-Trie is the first framework to achieve both simultaneously, and the Semantic Irrelevance Ratio introduced in Chapter 3 is the first metric designed to measure this property directly.

2.9.4 Concurrent and Subsequent Work: Positioning DCA-Trie

The research landscape has evolved rapidly since the foundational works of GCR (Luo *et al.*, 2024) and DoG (Li *et al.*, 2024). Most notably, **RouterKGQA** (Yuan *et al.*, 2026) proposes a specialized–general model routing framework that achieves strong performance with minimal inference cost. RouterKGQA employs constraint-aware reasoning paths and uses SBERT (nomic-embed-text-v1) for semantic filtering during answer selection, achieving 3.57 F1 point improvements over GCR on WebQSP and CWQ while requiring only 1.15 LLM calls per question (Yuan *et al.*, 2026).

Differentiation of DCA-Trie

While RouterKGQA and DCA-Trie both employ semantic relevance scoring, they operate at fundamentally different architectural levels:

Table 2.4 Architectural comparison of RouterKGQA and DCA-Trie.

Aspect	**RouterKGQA**	**DCA-Trie**
Intervention Point	Answer filtering (post-generation)	Token constraint (during generation)
Mechanism	Routes between specialized and general models	Modifies the constraint oracle directly
Semantic Scoring	Filters final candidate answers	Filters valid token sets at each decoding step
Faithfulness Guarantee	Probabilistic (repair when needed)	Deterministic (hard constraint at every step)
Metric Innovation	Efficiency (LLM calls)	Constraint quality (SIR)

RouterKGQA's semantic filtering operates on complete reasoning paths after generation, whereas DCA-Trie's semantic scoring constrains the *decoding process itself*–ensuring that the LLM never sees semantically irrelevant tokens during generation. This distinction is critical: DCA-Trie reduces the cognitive load on the LLM by narrowing the valid vocabulary at each step, while RouterKGQA relies on the LLM to generate paths then filters retrospectively.

Furthermore, DCA-Trie introduces the **Semantic Irrelevance Ratio (SIR)**, the first metric to isolate and quantify constraint permissiveness independently of downstream generation accuracy. RouterKGQA does not provide a direct measure of how many structurally valid but semantically irrelevant paths enter its candidate set; DCA-Trie's SIR metric fills this evaluation gap.

2.10 Synthesis of the Literature: Three Unresolved Gaps and the DCA-Trie Contribution

This section synthesizes the findings of the literature review by formalizing three distinct research gaps. Each limitation is then explicitly mapped to a proposed component of the DCA-Trie architecture and its corresponding project objective, establishing the foundation for the methodology detailed in Chapter 3.

2.10.1 Three Unresolved Gaps at the Intersection of Structural Faithfulness, Semantic Relevance, and Constraint Measurability

Gap 1: The Permissiveness Problem. GCR and DoG construct constraint oracles as $f(\mathcal{G}, \mathcal{E}_q)$, independent of question semantics and partial generation. For a three-hop CWQ question, Equation (2.4) shows this produces a valid token set containing up to 24,000 structurally correct but semantically irrelevant paths, forcing the LLM to use hallucination-prone parametric knowledge for discrimination, precisely the mechanism constrained decoding was designed to suppress (Luo *et al.*, 2024; Li *et al.*, 2024).

Gap 2: The Static-Dynamic Efficiency Tradeoff. GCR's precomputed static trie is computationally efficient but semantically blind. DoG's step-wise topological expansion introduces dynamism at the cost of per-step trie reconstruction and without semantic filtering (Li *et al.*, 2024; Willard and Louf, 2023). No existing framework achieves semantic responsiveness while preserving the efficiency of a precomputed oracle.

Gap 3: The Absence of a Permissiveness Metric. No existing work quantifies constraint permissiveness independently of downstream answer accuracy, making it impossible to measure progress on reducing irrelevant paths in the valid set directly or to compare frameworks on this property (Luo *et al.*, 2024; Li *et al.*, 2024).

Table 2.5 Consolidated gap matrix for KG-constrained decoding frameworks.

Framework	Faithfulness	Sem. Cond.	Dynamic	Precomp.	SIR
GCR (Luo *et al.*, 2024)	✓	No	No	✓	No
DoG (Li *et al.*, 2024)	✓	No	✓	No	No
RouterKGQA (Yuan *et al.*, 2026)	∼[a]	✓(answer)	✓	N/A	No
DCA-Trie v1 (proposed)	✓	✓(token)	No	✓	✓
DCA-Trie v2 (proposed)	✓	✓(token)	✓	Partial	✓

[a] Probabilistic faithfulness via repair mechanism; not guaranteed.

2.10.2 DCA-Trie as the Unified Response: Gap-to-Objective Mapping

Each gap maps directly to a DCA-Trie component and project objective. **Gap 1** is addressed by DCA-Trie's semantic relevance scorer applied as a hard admission criterion, formally upgrading the constraint oracle from $f(\mathcal{G}, \mathcal{E}_q)$ to $f(\mathcal{G}, \mathcal{E}_q, q, y_{<t})$, targeting project Objectives ii, iii, and iv. **Gap 2** is addressed by the two-variant design: DCA-Trie v1 applies semantic filtering at trie construction time, preserving precomputation efficiency (Objective iii); DCA-Trie v2 applies step-wise dynamic expansion with semantic filtering for maximum contextual responsiveness (Objective iv). **Gap 3** is addressed by the Semantic Irrelevance Ratio (SIR), a new metric introduced in Chapter 3 that quantifies the proportion of the valid set consisting of semantically irrelevant paths (Objective i).

DCA-Trie is the first constrained decoding framework to simultaneously address all three gaps, providing structural faithfulness, semantic relevance conditioning, and a dedicated permissiveness metric. RouterKGQA addresses efficiency through routing but does not provide deterministic faithfulness guarantees or constraint-level evaluation; DCA-Trie fills these critical gaps. These constitute the three novel contributions of this dissertation.

Table 2.6 DCA-Trie contribution mapping.

Research Gap	**DCA-Trie Component**	**Objective**	**Evaluated In**
Gap 1: Permissiveness Problem	Semantic scorer + threshold τ	Obj. ii, iii, iv	Ch. 3–4
Gap 2: Static-Dynamic Tradeoff	v1 (static) and v2 (dynamic) variants	Obj. iii, iv	Ch. 3–4
Gap 3: No Permissiveness Metric	Semantic Irrelevance Ratio (SIR)	Obj. i	Ch. 3–4

REFERENCES

Achiam, J., Adler, S., Agarwal, S., Ahmad, L., Akkaya, I., Aleman, F. L., *et al.* (2023). "GPT-4 Technical Report". In: *arXiv preprint arXiv:2303.08774*.

Agrawal, G., Kumarage, T., Alghamdi, Z., and Liu, H. (2024a). "Mindful-RAG: A Study of Points of Failure in Retrieval Augmented Generation". In: *arXiv preprint arXiv:2407.12216*.

Agrawal, G., Kumarage, T., Alghamdi, Z., and Liu, H. (2024b). "Mindful-RAG: A Study of Points of Failure in Retrieval-Augmented Generation". In: *arXiv preprint arXiv:2407.12216*.

Asai, A., Wu, Z., Wang, Y., Sil, A., and Hajishirzi, H. (2024). "Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection". In: *Proceedings of the Twelfth International Conference on Learning Representations (ICLR 2024)*.

Bollacker, K., Evans, C., Paritosh, P., Sturge, T., and Taylor, J. (2008). "Freebase: A Collaboratively Created Graph Database for Structuring Human Knowledge". In: *Proceedings of the 2008 ACM SIGMOD International Conference on Management of Data*, pp. 1247–1250.

Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., Agarwal, S., Herbert-Voss, A., Krueger, G., Henighan, T., Child, R., Ramesh, A., Ziegler, D., Wu, J., Winter, C., Hesse, C., Chen, M., Sigler, E., Litwin, M., Gray, S., Chess, B., Clark, J., Berner, C., McCandlish, S., Radford, A., Sutskever, I., and Amodei, D. (2020). "Language Models are Few-Shot Learners". In: *Advances in Neural Information Processing Systems*. Vol. 33. Curran Associates, Inc., pp. 1877–1901.

Chen, J., Xiao, S., Zhang, P., Luo, K., Lian, D., and Liu, Z. (2024). "BGE M3-Embedding: Multi-Lingual, Multi-Functionality, Multi-Granularity Text Embeddings Through Self-Knowledge Distillation". In: *arXiv preprint arXiv:2402.03216*.

De Cao, N., Izacard, G., Riedel, S., and Petroni, F. (2021). "Autoregressive Entity Retrieval". In: *arXiv preprint arXiv:2010.00904*. ICLR 2021.

Edge, D., Trinh, H., Cheng, N., Bradley, J., Chao, A., Mody, A., Truitt, S., and Larson, J. (2024). "From Local to Global: A Graph RAG Approach to Query-Focused Summarization". In: *arXiv preprint arXiv:2404.16130*.

Fredkin, E. (1960). "Trie Memory". In: *Communications of the ACM* 3.9, pp. 490–499.

Geng, S., Josifoski, M., Peyrard, M., and West, R. (2023). *Grammar-Constrained Decoding for Structured NLP Tasks without Finetuning*. arXiv:2305.13971.

Grattafiori, A. *et al.* (2024). "The Llama 3 Herd of Models". In: arXiv: 2407.21783 [cs.AI]. URL: https://arxiv.org/abs/2407.21783.

He, G., Lan, Y., Jiang, J., Zhao, W. X., and Wen, J.-R. (2021a). "Improving Multi-Hop Knowledge Base Question Answering by Learning Intermediate Supervision Signals". In: *Proceedings of the 14th ACM International Conference on Web Search and Data Mining*, pp. 553–561.

— (2021b). "Improving Multi-Hop Knowledge Base Question Answering by Learning Intermediate Supervision Signals". In: *Proceedings of the 14th ACM International Conference on Web Search and Data Mining (WSDM)*. ACM, pp. 553–561.

Holtzman, A., Buys, J., Du, L., Forbes, M., and Choi, Y. (2020). "The Curious Case of Neural Text Degeneration". In: *International Conference on Learning Representations*.

Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, W., Feng, X., Qin, B., and Liu, T. (2023). "A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions". In: *arXiv preprint arXiv:2311.05232*.

Izacard, G. and Grave, E. (2021). "Leveraging Passage Retrieval with Generative Models for Open Domain Question Answering". In: *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics*, pp. 874–880.

Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y., Madotto, A., and Fung, P. (2023). "Survey of Hallucination in Natural Language Generation". In: *ACM Computing Surveys* 55.12. arXiv:2202.03629, pp. 1–38.

Jiang, Z., Xu, F. F., Gao, L., Sun, Z., Liu, Q., Dwivedi-Yu, J., Yang, Y., Callan, J., and Neubig, G. (2023). "Active Retrieval Augmented Generation". In: *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, pp. 7969–7992.

Kandpal, N., Deng, H., Roberts, A., Wallace, E., and Raffel, C. (2023). "Large Language Models Struggle to Learn Long-Tail Knowledge". In: *Proceedings of the 40th International Conference on Machine Learning (ICML)*. PMLR, pp. 15696–15707.

Karpukhin, V., Oğuz, B., Min, S., Lewis, P., Wu, L., Edunov, S., Chen, D., and Yih, W.-t. (2020). "Dense Passage Retrieval for Open-Domain Question Answering". In: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, pp. 6769–6781.

Lan, Y., He, G., Jiang, J., Jiang, J., Zhao, W. X., and Wen, J.-R. (2022). "Complex Knowledge Base Question Answering: A Survey". In: *arXiv preprint arXiv:2108.06688*.

Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.-t., Rocktäschel, T., Riedel, S., and Kiela, D. (2020). "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks". In: *Advances in Neural Information Processing Systems*.

Li, K., Zhang, T., Wu, X., Luo, H., Glass, J., and Meng, H. (2024). *Decoding on Graphs: Faithful and Sound Reasoning on Knowledge Graphs through Generation of Well-Formed Chains*. arXiv:2410.18415.

Lin, S., Hilton, J., and Evans, O. (2022). "TruthfulQA: Measuring How Models Mimic Human Falsehoods". In: *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*. Association for Computational Linguistics, pp. 3214–3252.

Luo, L., Tang, Z., Shi, Z., Wei, X., Xu, Y., Zhang, W., Liu, Z., and Tang, J. (2024). *Graph-Constrained Reasoning: Faithful Reasoning on Knowledge Graphs for Large Language Models*. arXiv:2410.13080.

Luo, L., Zhao, Z., Haffari, G., Li, Y.-F., Gong, C., and Pan, S. (2025). "Graph-constrained Reasoning: Faithful Reasoning on Knowledge Graphs with Large Language Models". In: *Forty-second International Conference on Machine Learning*. URL: https://openreview.net/forum?id=Fr7kH2SFq7.

Mallen, A., Asai, A., Zhong, V., Das, R., Khashabi, D., and Hajishirzi, H. (2023). "When Not to Trust Language Models: Investigating Effectiveness of Parametric and Non-Parametric Memories". In: *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*. Association for Computational Linguistics, pp. 9802–9822.

OpenAI (2024). *GPT-4 Technical Report*. arXiv: 2303.08774 [cs.CL]. URL: https://arxiv.org/abs/2303.08774.

Perez, E., Ringer, S., Lukošiūtė, K., Nguyen, K., Chen, E., Heiner, S., Pettit, C., Olsson, C., Kundu, S., Kadavath, S., Jones, A., Chen, A., Mann, B., Israel, B., Bowman, S. R., Clark, J., Ganguli, D., Askell, A., and Hernandez, D. (2023). "Discovering Language Model Behaviors with Model-Written Evaluations". In: *Findings of the Association for Computational Linguistics: ACL 2023*. Association for Computational Linguistics, pp. 13387–13434.

Reimers, N. and Gurevych, I. (2019). "Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks". In: *arXiv preprint arXiv:1908.10084*. EMNLP 2019.

Saxena, A., Chakrabarti, S., and Talukdar, P. (2022). "Sequence-to-Sequence Knowledge Graph Completion and Question Answering". In: *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*. Association for Computational Linguistics, pp. 2814–2828.

Shi, J., Cao, S., Hou, L., Li, J., and Zhang, H. (2021). "transferNet: An Effective and Transparent Framework for Multi-Hop Question Answering over Relation Graph". In: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, pp. 4149–4158.

Talmor, A. and Berant, J. (2018). "The Web as a Knowledge-Base for Answering Complex Questions". In: *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics*.

Touvron, H., Martin, L., Stone, K., *et al.* (2023). "LLaMA: Open and Efficient Foundation Language Models". In: *arXiv preprint arXiv:2302.13971*.

Trivedi, H., Balasubramanian, N., Khot, T., and Sabharwal, A. (2023). "Interleaving Retrieval with Chain-of-Thought Reasoning for Knowledge-Intensive Multi-Step Questions". In: *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*. Association for Computational Linguistics, pp. 10014–10037.

Wang, X., Wei, J., Schuurmans, D., Le, Q. V., Chi, E. H., Narang, S., Chowdhery, A., and Zhou, D. (2023). "Self-Consistency Improves Chain of Thought Reasoning in Language Models". In: *The Eleventh International Conference on Learning Representations*. URL: `https://openreview.net/forum?id=1PL1NIMMrw`.

Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., Chi, E., Le, Q. V., and Zhou, D. (2022). "Chain-of-Thought Prompting Elicits Reasoning in Large Language

Models". In: *Advances in Neural Information Processing Systems*. Vol. 35. Curran Associates, Inc., pp. 24824–24837.

Willard, B. T. and Louf, R. (2023). *Efficient Guided Generation for Large Language Models*. arXiv:2307.09702.

Xu, Z., Jain, S., and Kankanhalli, M. (2025). *Hallucination is Inevitable: An Innate Limitation of Large Language Models*. arXiv: 2401.11817 [cs.CL]. URL: https://arxiv.org/abs/2401.11817.

Yao, S., Yu, D., Zhao, J., Shafran, I., Griffiths, T., Cao, Y., and Narasimhan, K. (2023). "Tree of Thoughts: Deliberate Problem Solving with Large Language Models". In: *Advances in Neural Information Processing Systems*. Vol. 36. Curran Associates, Inc.

Yasunaga, M., Ren, H., Bosselut, A., Liang, P., and Leskovec, J. (2021). "QA-GNN: Reasoning with Language Models and Knowledge Graphs for Question Answering". In: *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*. Association for Computational Linguistics, pp. 535–546.

Yih, W.-t., Richardson, M., Meek, C., Chang, M.-W., and Suh, J. (2016). "The Value of Semantic Parse Labeling for Knowledge Base Question Answering". In: *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics*, pp. 201–206.

Yuan, B., Deng, H., Liu, X., and Zhang, M. (2026). "RouterKGQA: Specialized–General Model Routing for Constraint-Aware Knowledge Graph Question Answering". In: arXiv: 2603.20017 [cs.CL]. URL: https://arxiv.org/abs/2603.20017.